

cs109a_finalproject_submission

December 7, 2017

0.1 Packages used that are not included in Jupyter Notebook.

Fuzzywuzzy-- string matching using Levenshtein Distance to calculate the differences between strings

<https://pypi.python.org/pypi/fuzzywuzzy>

Geopy-- locates the coordinates of addresses, cities, countries, and landmarks across the globe using third-party geocoders and other data sources

<https://pypi.python.org/pypi/geopy>

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import pickle
import seaborn as sns
import requests
from bs4 import BeautifulSoup
from fuzzywuzzy import fuzz
import cartopy.crs as ccrs
import matplotlib.pyplot as plt
import matplotlib.patches as mpatches
import matplotlib.pyplot as plt
import shapely.geometry as sgeom
import cartopy.io.shapereader as shpreader
import cartopy
import pickle
from geopy.geocoders import Nominatim
import math
import plotly.plotly as py
import plotly

from sklearn.preprocessing import MinMaxScaler
from sklearn.linear_model import LassoCV
from sklearn.linear_model import RidgeCV
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestRegressor
```

```

from sklearn.preprocessing import PolynomialFeatures
from sklearn.decomposition import PCA
from sklearn.ensemble import RandomForestRegressor
% matplotlib inline

```

/Users/jeanettejin/anaconda/lib/python3.6/site-packages/fuzzywuzzy/fuzz.py:35: UserWarning: Using slow pure-python SequenceMatcher. Install python-Levenshtein to remove these warnings.
 warnings.warn('Using slow pure-python SequenceMatcher. Install python-Levenshtein to remove these warnings.')

1 Data Collection

1.1 FBI Data

Here, we scrape from the FBI webpage, getting the hyperlinks for each year, and then we navigate to violent crimes and murders, getting the URL for each year's MSA table.

In [2]: *# Code that scrapes data from each webpage for each year*

```

## for each year
req = requests.get("https://ucr.fbi.gov/ucr-publications")
page = req.text
soup = BeautifulSoup(page, 'html.parser')

# scrape links from main website
links = []
span = soup.find_all('span', class_='castle-body')[4]
a = span.find_all('a')

for link in a:
    links.append(link['href'])
links = links[0:11]
links

# redefine link for 09
links[7] = 'https://www2.fbi.gov/ucr/cius2009/index.html'

## going to violent crimes
violence = []

for link in links:
    req = requests.get(link)
    page = req.text
    soup = BeautifulSoup(page, 'html.parser')
    href = soup.find_all('a', href = True, text = 'Violent Crime')[0]['href']
    if 'https' in href:
        violence.append(href)

    else:

```

```

        if 'index.html' in link:
            link = link.replace('index.html', '')
            violence.append(link + href)
        else:
            violence.append(link + '/' + href)

## going to MSA table
table = []
for v in violence:
    page = requests.get(v).text
    soup = BeautifulSoup(page, 'html.parser')
    a = soup.find_all('a', href = True)
    for i in a:
        if 'Statistical Areas' in i.text:
            if 'https' in i['href']:
                table.append(i['href'])
            else:
                if 'violent_crime' in v:
                    v = v.replace('offenses/violent_crime/index.html', '')
                    table.append(v + 'data/table_06.html')

table

```

```

Out[2]: ['https://ucr.fbi.gov/crime-in-the-u.s/2016/crime-in-the-u.s.-2016/tables/table-4',
'https://ucr.fbi.gov/crime-in-the-u.s/2015/crime-in-the-u.s.-2015/tables/table-6',
'https://ucr.fbi.gov/crime-in-the-u.s/2014/crime-in-the-u.s.-2014/tables/table-6',
'https://ucr.fbi.gov/crime-in-the-u.s/2013/crime-in-the-u.s.-2013/tables/6tabledatadecp',
'https://ucr.fbi.gov/crime-in-the-u.s/2012/crime-in-the-u.s.-2012/tables/6tabledatadecp',
'https://ucr.fbi.gov/crime-in-the-u.s/2011/crime-in-the-u.s.-2011/tables/table-6',
'https://ucr.fbi.gov/crime-in-the-u.s/2010/crime-in-the-u.s.-2010/tables/table-6',
'https://www2.fbi.gov/ucr/cius2009/data/table_06.html',
'http://www2.fbi.gov/ucr/cius2008/data/table_06.html',
'http://www2.fbi.gov/ucr/cius2007/data/table_06.html',
'http://www2.fbi.gov/ucr/cius2006/data/table_06.html']

```

We define a function to parse the tables for each year. This function was adapted from <http://srome.github.io/Parsing-HTML-Tables-in-Python-with-BeautifulSoup-and-pandas/>, but makes many changes to td/th tag checks, dataframe indexing, changes to find the dimensions of the dataframe, and changes to float/int recast.

```

In [3]: # Adapted from http://srome.github.io/Parsing-HTML-Tables-in-Python-with-BeautifulSoup-a
        # changes to td/th checks
        # changes to dataframe indexing
        # changes to how we find dim of dataframe
        # changes to float/int recast

def parse_table(table):
    n_columns = 0

```

```

n_rows = 0
column_names = []

# Find number of rows and columns
# we also find the column titles if we can
colidx = 0
rows = table.find_all('tr')
n_rows = len(rows)-1
td_tags = rows[0].find_all('td')
th_tags = rows[0].find_all('th')
n_columns = len(th_tags) + len(td_tags)
for th in th_tags:
    column_names.append(th.get_text().strip())

# Safeguard on Column Titles
if len(column_names) > 0 and len(column_names) != n_columns:
    raise Exception("Column titles do not match the number of columns")

columns = column_names if len(column_names) > 0 else range(0,n_columns)
df = pd.DataFrame(columns = columns,
                  index= range(0,n_rows))
row_marker = 0

# find all rows and iterate over each row (skip first row- those are the columns)
for row in rows[1:]:
    # check if 2 th's (means at Abilene)
    numb_cols = len(row.find_all('td')) + len(row.find_all('th'))
    if numb_cols == n_columns:
        column_marker = 0
    else:
        df.iat[row_marker,0] = ''
        column_marker = 1

# find each column header (there are 2 column headers and 10 td's)
colheads = row.find_all('th')
if row_marker == 0:
    print(df.iat[0,0])
for colhead in colheads:
    df.iat[row_marker,column_marker] = colhead.get_text().strip()
    column_marker+=1

# find the 10 td's and append to dataframe
columns = row.find_all('td')
for column in columns:
    df.iat[row_marker,column_marker] = column.get_text().strip()
    column_marker += 1
if len(columns) > 0:
    row_marker += 1

```

```
return df
```

We use the function defined above to scrape all elements in the table and store them in a dictionary of dataframes with keys = year

```
In [4]: # store the tables in a dictionary with key = year
        yeartbl = {}
        year = 2016
        for t in table:
            page = requests.get(t).text
            soup = BeautifulSoup(page, 'lxml')
            i = 0
            # pick the first table with many rows
            while len(soup.find_all('table')[i].find_all('tr')) < 200:
                i += 1
            a = soup.find_all('table')[i]
            tempdf = parse_table(a)
            yeartbl[year] = tempdf
            year = year - 1
```

```
nan
nan
nan
nan
nan
nan
nan
nan
nan
nan
nan
nan
```

An example:

```
In [5]: yeartbl[2006].head()
```

```
Out[5]: Metropolitan Statistical Area \
0      Abilene, TX M.S.A.1
1
2
3
4
```

	Counties/principal cities	Population	Violent crime \
0		162,776	
1	Includes Callahan,1 Jones, and Taylor Counties		
2	City of Abilene	118,009	554

3	Total area actually reporting	100.0%	638
4	Rate per 100,000 inhabitants		391.9

	Murder and nonnegligent manslaughter	Forcible rape	Robbery	\
0				
1				
2	5	67	107	
3	6	75	109	
4	3.7	46.1	67.0	

	Aggravated assault	Property crime	Burglary	Larceny-theft	Motor vehicle theft
0					
1					
2	375	5,045	1,282	3,460	303
3	448	5,741	1,531	3,852	358
4	275.2	3,526.9	940.6	2,366.4	219.9

Changing column names:

```
In [6]: # replace old column names with easier to use
new_col_names = ['MSA', 'Counties', 'Population', 'Violent_Crime', 'Murder', 'Rape', 'Ro
            'Vehicle_Theft']

i = 2016

while i != 2005:
    yeartbl[i].columns = new_col_names
    i = i - 1
```

We are only interested in the MSA and its murder counts for each year. For each table in our dictionary we find this information and put them into another dictionary of dataframe with key = year

```
In [7]: # since we only care about murder at the moment lets change the type to numeric
i = 2016

murder_stats = {}
while i != 2005:
    temp = yeartbl[i]

    # remove commas to change type
    temp['Murder'] = temp['Murder'].str.replace(",", "")
    temp['Murder'] = pd.to_numeric(temp['Murder'])
    temp['MSA'][temp['MSA'] == ''] = np.NaN
    temp['MSA'] = temp['MSA'].fillna(method = 'ffill')
    murders = temp[temp["Counties"].str.contains("Total area actually reporting")][['MSA
    murder_stats[i] = murders
    i = i - 1
```

An example:

```
In [8]: murder_stats[2006].head()
```

```
Out[8]:
```

		MSA	Murder
3	Abilene, TX	M.S.A.1	6.0
8	Albany, GA	M.S.A.	11.0
16	Albany-Schenectady-Troy, NY	M.S.A.	18.0
21	Albuquerque, NM	M.S.A.	72.0
27	Alexandria, LA	M.S.A.	10.0

```
In [9]: # lengths of each dataframe
        for i in range(2006,2017):
            print(len(murder_stats[i]))
```

```
357
364
357
378
368
376
389
384
377
377
386
```

1.2 Census Data

We download csv files of for year 2006 - 2016 for the following features take from the Census Data by each MSA

S0101: Age and Sex S1201: Marital Status S1501: Educational Attainment S1701: Poverty Status in the Past 12 Months S1903: Median income in the past 12 months B02001: Race B01003: Total Population

Below we define some functions to help us with the cleaning process:

```
In [10]: # makes the df column names the first row
        def rename(df):
            header = df.iloc[0]
            df = df.iloc[1:]
            df = df.rename(columns = header)
            return df

        # names of the csv files
        def makefile_names(code, some_list):
            for i in list(range(6, 17)):
                num = "%02d" % (i,)
```

```

        if i == 6:
            some_list.append('Census Data/ACS_06_EST_' + code + '_with_ann.csv')
        else:
            some_list.append('Census Data/ACS_' + num + '_1YR_' + code + '_with_an
    return some_list

In [11]: # the codes of each feature
codes = ['S0101', 'S1201', 'S1501', 'S1701', 'S1903', 'B02001', 'B01003']

filenames = {}

names = ['age_sex', 'marital', 'education', 'poverty', 'income', 'race', 'population']

for i in range(0, len(names)):
    list_name = []
    filenames[names[i]] = makefile_names(codes[i], list_name)

```

1.2.1 Feature: Population

We use the first feature population to put each MSA into a dictionary with key = unique identifier. For each year we iterate through the dataframe's MSA and check to see if it exists within our dictionary. If not, we add it to our dictionary. This ensures that we will have nan values for MSAs that do not appear throughout all the years.

```

In [12]: dict_pop_year = {}

# make a dict of dict with each being {msa_code: population_year}
for i in range(0, len(filenames['population'])):

    # load as df
    df_pop_year = pd.read_csv(filenames['population'][i], encoding = "ISO-8859-1")
    print(len(df_pop_year))

    # clean
    df_pop_year = rename(df_pop_year)
    df_pop_year[['Id2']] = df_pop_year[['Id2']].astype(int)
    df_pop_year['Geography'] = df_pop_year['Geography'].str.replace(' Metro Area', '')

    # make a dict of id and name of msa
    if i == 0:
        msa_dict = dict(zip(df_pop_year['Id2'], df_pop_year['Geography']))

    else:
        for code,name in zip(df_pop_year['Id2'], df_pop_year['Geography']):
            msa_dict.setdefault(code, name)

dict_pop = dict(zip(df_pop_year['Id2'], df_pop_year['Estimate; Total']))

```



```

dict_pop_year['pop' + "_%02d" % (i + 6,)] = dict_pop

# printed are the number of unique msa's per year!

368
370
370
375
375
375
375
389
389
389
390

```

We use the dictionary with key = Unique identifier and value the names of the MSA and create a df with all of the populations for each year.

```

In [13]: # make dataframe using dict keys
msa_df = pd.DataFrame.from_dict(msa_dict, orient = 'index')
msa_df.columns = ['msa_name']

# add to dict using keys
for key in dict_pop_year:
    msa_df[key] = pd.Series(dict_pop_year[key])

# show
msa_df.head()

```

```

Out[13]:

```

	msa_name	pop_06	pop_07	pop_08	pop_09	\
10180	Abilene, TX	158548	159439	160012	160266	
10380	Aguadilla-Isabela-San Sebastián, PR	336502	335201	339193	342495	
10420	Akron, OH	700943	699356	698553	699935	
10500	Albany, GA	165062	162767	163074	164238	
10580	Albany-Schenectady-Troy, NY	850957	853358	853919	857592	

	pop_10	pop_11	pop_12	pop_13	pop_14	pop_15	pop_16
10180	164941	165858	167800	168144	166900	168922	170860
10380	305988	304633	304727	327847	322079	313209	309764
10420	702951	701456	702262	705686	703825	704243	702221
10500	162659	161617	155019	156277	152596	156997	152506
10580	870832	871478	874646	877905	880167	881830	881839

We do this for every feature using the unique identifier to add features to our dataframe

1.2.2 Feature: Sex and Year

```
In [14]: dict_agesex_year = {}
```

```
for i in range(0, len(filenamees['age_sex'])):

    # load and clean
    df_agesex_year = pd.read_csv(filenamees['age_sex'][i], encoding = 'ISO-8859-1')
    df_agesex_year = rename(df_agesex_year)

    # get columns names
    gender_columns = ['Male; Estimate; Total population', 'Female; Estimate; Total popu
    age_columns = [x for x in df_agesex_year.columns.values.tolist() if x.startswith("T
    if len(age_columns) == 0:
        age_columns = df_agesex_year.columns[df_agesex_year.columns.str.contains('Total
    int_columns = gender_columns + ['Total; Estimate; Total population', 'Id2']

    # change type
    df_agesex_year[int_columns] = df_agesex_year[int_columns].astype(int)
    df_agesex_year[age_columns] = df_agesex_year[age_columns].astype(float)

    # make a column for the percent of male and females in each msa
    df_agesex_year['Percent_male'] = df_agesex_year[gender_columns[0]] / df_agesex_year
    df_agesex_year['Percent_female'] = df_agesex_year[gender_columns[1]] / df_agesex_ye

    # select columns
    add_columns_agesex = ['Percent_male', 'Percent_female'] + age_columns

    # fill dict
    for feature in add_columns_agesex:
        dict_agesex = dict(zip(df_agesex_year['Id2'], df_agesex_year[feature]))
        dict_agesex_year[feature + "_%02d" % (i + 6,)] = dict_agesex

    # add to dataframe
    for key in dict_agesex_year:
        msa_df[key] = pd.Series(dict_agesex_year[key])

    # rename and show
    msa_df.columns = msa_df.columns.str.replace('Total; Estimate; ', '')
    msa_df.columns = msa_df.columns.str.replace('Total population - ', '')
    msa_df.head()
```

```
Out[14]:
```

	msa_name	pop_06	pop_07	pop_08	pop_09	\
10180	Abilene, TX	158548	159439	160012	160266	
10380	Aguadilla-Isabela-San Sebastián, PR	336502	335201	339193	342495	
10420	Akron, OH	700943	699356	698553	699935	
10500	Albany, GA	165062	162767	163074	164238	
10580	Albany-Schenectady-Troy, NY	850957	853358	853919	857592	

	pop_10	pop_11	pop_12	pop_13	pop_14	...	\
10180	164941	165858	167800	168144	166900	...	
10380	305988	304633	304727	327847	322079	...	
10420	702951	701456	702262	705686	703825	...	
10500	162659	161617	155019	156277	152596	...	
10580	870832	871478	874646	877905	880167	...	

	AGE - 40 to 44 years_16	AGE - 45 to 49 years_16	\
10180	4.4	5.3	
10380	6.0	6.4	
10420	5.5	6.5	
10500	5.4	5.6	
10580	5.6	6.7	

	AGE - 50 to 54 years_16	AGE - 55 to 59 years_16	\
10180	6.2	5.2	
10380	6.5	6.5	
10420	7.3	7.4	
10500	6.1	6.6	
10580	7.3	7.5	

	AGE - 60 to 64 years_16	AGE - 65 to 69 years_16	\
10180	5.9	5.3	
10380	6.6	6.4	
10420	6.9	5.7	
10500	6.6	5.6	
10580	6.3	5.6	

	AGE - 70 to 74 years_16	AGE - 75 to 79 years_16	\
10180	3.3	2.7	
10380	5.0	3.5	
10420	3.9	2.8	
10500	3.3	2.4	
10580	3.8	2.2	

	AGE - 80 to 84 years_16	AGE - 85 years and over_16
10180	2.1	1.8
10380	2.1	1.8
10420	1.9	2.4
10500	2.0	1.5
10580	2.0	2.6

[5 rows x 232 columns]

1.2.3 Feature: Marriage Status

In [15]: dict_marital_year = {}

```

for i in range(0, len(filenamees['marital'])):

    # load and clean
    df_marital = pd.read_csv(filenamees['marital'][i], encoding='cp1252')
    df_marital = rename(df_marital)
    df_marital = df_marital.loc[:,~df_marital.columns.duplicated()]
    df_marital = df_marital.replace("N", np.nan)

    # features
    status_columns = ['Now married (except separated); Estimate; Population 15 years and over']

    # change type
    df_marital[status_columns] = df_marital[status_columns].astype(float)
    df_marital['Id2'] = df_marital['Id2'].astype(float)

    # fill dict with dicts
    for feature in status_columns:
        dict_marital = dict(zip(df_marital['Id2'], df_marital[feature]))
        dict_marital_year[feature + "_%02d" % (i + 6,)] = dict_marital

    # add to df
    for key in dict_marital_year:
        msa_df[key] = pd.Series(dict_marital_year[key])

    # clean column names and show
    msa_df.columns = msa_df.columns.str.replace('; Estimate; Population 15 years and over', '')
    msa_df.columns = msa_df.columns.str.replace('except separated', '')
    msa_df.head()

```

```

Out[15]:

```

	msa_name	pop_06	pop_07	pop_08	pop_09	\
10180	Abilene, TX	158548	159439	160012	160266	
10380	Aguadilla-Isabela-San Sebastián, PR	336502	335201	339193	342495	
10420	Akron, OH	700943	699356	698553	699935	
10500	Albany, GA	165062	162767	163074	164238	
10580	Albany-Schenectady-Troy, NY	850957	853358	853919	857592	

	pop_10	pop_11	pop_12	pop_13	pop_14	...	\
10180	164941	165858	167800	168144	166900	...	
10380	305988	304633	304727	327847	322079	...	
10420	702951	701456	702262	705686	703825	...	
10500	162659	161617	155019	156277	152596	...	
10580	870832	871478	874646	877905	880167	...	

	Now married	()_15	Widowed_15	Divorced_15	Separated_15	\
10180		45.4	6.6	12.0	3.3	

10380	41.7	7.1	12.0	3.8
10420	46.7	6.3	11.9	1.3
10500	39.4	7.8	12.6	5.2
10580	45.1	5.9	9.8	1.9

	Never married_15	Now married ()_16	Widowed_16	Divorced_16	\
10180	32.8	46.8	5.5	10.9	
10380	35.3	39.5	7.2	12.6	
10420	33.8	46.9	5.9	11.5	
10500	35.0	39.8	6.8	12.0	
10580	37.2	44.2	6.5	10.2	

	Separated_16	Never married_16
10180	2.5	34.2
10380	3.1	37.7
10420	1.6	34.1
10500	3.3	38.2
10580	2.2	36.9

[5 rows x 287 columns]

1.2.4 Feature: Education

```
In [16]: dict_education_year = {}
```

```
for i in range(0, len(filenamees['education'])):

    # load and clean
    df_education = pd.read_csv(filenamees['education'][i], encoding='cp1252')
    df_education = rename(df_education)

    # features (gov naming conventions suck)
    if i < 4:
        education_columns = ['Total; Estimate; Population 25 years and over - Percent h

    if 4 <= i <= 8:
        education_columns = ['Total; Estimate; Percent high school graduate or higher',

    if i > 8:
        education_columns = ['Percent; Estimate; Percent high school graduate or higher

    # change datatype
    df_education[education_columns] = df_education[education_columns].astype(float)
    df_education['Id2'] = df_education['Id2'].astype(int)

    # create dict of dicts
    for feature in education_columns:
```

```

dict_education = dict(zip(df_education['Id2'], df_education[feature]))
dict_education_year[feature + "_%02d" % (i + 6,) ] = dict_education

# add to dataframe
for key in dict_education_year:
    msa_df[key] = pd.Series(dict_education_year[key])

# clean column names and show
msa_df.columns = msa_df.columns.str.replace('Total; Estimate;', '')
msa_df.columns = msa_df.columns.str.replace('Percent; Estimate;', '')
msa_df.columns = msa_df.columns.str.replace('Population 25 years and over - ', '')
msa_df.head()

```

```

Out[16]:

```

	msa_name	pop_06	pop_07	pop_08	pop_09	\
10180	Abilene, TX	158548	159439	160012	160266	
10380	Aguadilla-Isabela-San Sebastián, PR	336502	335201	339193	342495	
10420	Akron, OH	700943	699356	698553	699935	
10500	Albany, GA	165062	162767	163074	164238	
10580	Albany-Schenectady-Troy, NY	850957	853358	853919	857592	

	pop_10	pop_11	pop_12	pop_13	pop_14	\
10180	164941	165858	167800	168144	166900	
10380	305988	304633	304727	327847	322079	
10420	702951	701456	702262	705686	703825	
10500	162659	161617	155019	156277	152596	
10580	870832	871478	874646	877905	880167	

	...	\
10180	...	
10380	...	
10420	...	
10500	...	
10580	...	

	Percent high school graduate or higher_12	\
10180	83.5	
10380	66.0	
10420	90.6	
10500	79.9	
10580	92.0	

	Percent bachelor's degree or higher_12	\
10180	21.2	
10380	20.5	
10420	28.4	
10500	18.4	
10580	34.1	

	Percent high school graduate or higher_13 \
10180	83.8
10380	63.3
10420	91.1
10500	80.1
10580	92.1

	Percent bachelor's degree or higher_13 \
10180	22.1
10380	18.2
10420	29.7
10500	16.4
10580	34.3

	Percent high school graduate or higher_14 \
10180	86.1
10380	67.6
10420	91.1
10500	81.5
10580	92.5

	Percent bachelor's degree or higher_14 \
10180	21.6
10380	19.8
10420	29.9
10500	19.0
10580	35.4

	Percent high school graduate or higher_15 \
10180	87.7
10380	68.7
10420	91.5
10500	82.3
10580	92.2

	Percent bachelor's degree or higher_15 \
10180	20.7
10380	20.4
10420	30.1
10500	21.1
10580	35.3

	Percent high school graduate or higher_16 \
10180	86.7
10380	67.3
10420	91.8
10500	84.8
10580	91.6

	Percent bachelor's degree or higher_16
10180	20.4
10380	17.8
10420	31.0
10500	18.6
10580	37.4

[5 rows x 309 columns]

1.2.5 Feature Poverty

```
In [17]: dict_poverty_year = {}
```

```
for i in range(0, len(filenamees['poverty'])):

    # load and clean
    df_poverty = pd.read_csv(filenamees['poverty'][i], encoding='cp1252')
    df_poverty = rename(df_poverty)

    # find column name and set (a difference of a space between percent and below...)
    a = 'Percent below poverty level; Estimate; Population for whom poverty status is
    b = 'Percent below poverty level; Estimate; Population for whom poverty status is d

    if a in df_poverty:
        poverty_column = a

    if b in df_poverty:
        poverty_column = b

    # change dtype
    df_poverty[poverty_column] = df_poverty[poverty_column].astype(float)
    df_poverty['Id2'] = df_poverty['Id2'].astype(int)

    # make a dict of dicts
    dict_poverty = dict(zip(df_poverty['Id2'], df_poverty[poverty_column]))
    dict_poverty_year['poverty' + "_%02d" % (i + 6,) ] = dict_poverty

    # add each dict in dict to msa df
    for key in dict_poverty_year:
        msa_df[key] = pd.Series(dict_poverty_year[key])

msa_df.head()
```

```
Out[17]:
```

	msa_name	pop_06	pop_07	pop_08	pop_09	\
10180	Abilene, TX	158548	159439	160012	160266	
10380	Aguadilla-Isabela-San Sebastián, PR	336502	335201	339193	342495	
10420	Akron, OH	700943	699356	698553	699935	

10500		Albany, GA	165062	162767	163074	164238
10580		Albany-Schenectady-Troy, NY	850957	853358	853919	857592

	pop_10	pop_11	pop_12	pop_13	pop_14	...	poverty_07	\
10180	164941	165858	167800	168144	166900	...	15.1	
10380	305988	304633	304727	327847	322079	...	57.6	
10420	702951	701456	702262	705686	703825	...	13.4	
10500	162659	161617	155019	156277	152596	...	21.5	
10580	870832	871478	874646	877905	880167	...	10.1	

	poverty_08	poverty_09	poverty_10	poverty_11	poverty_12	poverty_13	\
10180	16.1	14.5	18.0	16.9	20.5	16.8	
10380	54.5	55.4	53.8	54.0	53.7	53.8	
10420	12.1	14.7	15.5	16.6	15.7	15.4	
10500	23.3	23.2	27.7	28.4	26.9	24.9	
10580	10.5	9.9	11.5	11.5	11.0	12.5	

	poverty_14	poverty_15	poverty_16
10180	16.5	13.9	17.2
10380	53.4	53.7	54.1
10420	13.3	14.2	13.8
10500	25.3	24.6	25.0
10580	11.7	10.2	9.9

[5 rows x 320 columns]

1.2.6 Feature: Income

```
In [18]: dict_income_year = {}
```

```
for i in range(0, len(filenamees['income'])):

    # load and clean
    df_income = pd.read_csv(filenamees['income'][i], encoding = 'cp1252')
    df_income = rename(df_income)

    # get column name
    income_column = [x for x in df_income.columns.values.tolist() if x.startswith('Medi

    # dtypes
    df_income[income_column] = df_income[income_column].astype(float)
    df_income['Id2'] = df_income['Id2'].astype(int)

    # dict of dicts
    dict_income = dict(zip(df_income['Id2'], df_income[income_column]))
    dict_income_year['median_income' + "_%02d" % (i + 6,) ] = dict_income

# add each dict in dict to msa df
```

```
for key in dict_income_year:
    msa_df[key] = pd.Series(dict_income_year[key])
```

```
msa_df.head()
```

```
Out[18]:
```

	msa_name	pop_06	pop_07	pop_08	pop_09	\
10180	Abilene, TX	158548	159439	160012	160266	
10380	Aguadilla-Isabela-San Sebastián, PR	336502	335201	339193	342495	
10420	Akron, OH	700943	699356	698553	699935	
10500	Albany, GA	165062	162767	163074	164238	
10580	Albany-Schenectady-Troy, NY	850957	853358	853919	857592	

	pop_10	pop_11	pop_12	pop_13	pop_14	...	\
10180	164941	165858	167800	168144	166900	...	
10380	305988	304633	304727	327847	322079	...	
10420	702951	701456	702262	705686	703825	...	
10500	162659	161617	155019	156277	152596	...	
10580	870832	871478	874646	877905	880167	...	

	median_income_07	median_income_08	median_income_09	median_income_10	\
10180	39369.0	41961.0	42931.0	40630.0	
10380	12139.0	13152.0	13470.0	14313.0	
10420	47898.0	50036.0	47482.0	46521.0	
10500	36402.0	38989.0	36218.0	34002.0	
10580	55129.0	58765.0	57677.0	55796.0	

	median_income_11	median_income_12	median_income_13	median_income_14	\
10180	40659.0	43407.0	44149.0	44303.0	
10380	14951.0	15339.0	15323.0	15886.0	
10420	47032.0	49731.0	49984.0	50538.0	
10500	32775.0	34469.0	34756.0	39071.0	
10580	58617.0	60625.0	59626.0	62265.0	

	median_income_15	median_income_16
10180	47420.0	48016.0
10380	14485.0	14546.0
10420	51580.0	51598.0
10500	40143.0	40667.0
10580	63080.0	65855.0

```
[5 rows x 331 columns]
```

1.2.7 Feature Race

```
In [19]: dict_race_year = {}
```

```
for i in range(0, len(filenamees['race'])):
```

```

# load and clean
df_race = pd.read_csv(filenamees['race'][i], encoding= 'cp1252')
df_race = rename(df_race)

# finding column names of interest
df_race.columns = df_race.columns.str.replace('Estimate;', '')
df_race.rename(columns={" Total:" : 'TOTAL'}, inplace = True)
df_race = df_race[df_race.columns.drop(list(df_race.filter(regex='Margin')))]
column_divide = df_race.columns.tolist()[3 :-2]
column_divide.append('Id2')

# change datatype
df_race[column_divide] = df_race[column_divide].astype(int)

names = []

# getting percentages
for feature in column_divide[1:-1]:
    df_race[feature + '_Percent'] = df_race[feature] / df_race['TOTAL']
    names.append(feature + '_Percent')

# dict of dicts
for feature in names:
    dict_race = dict(zip(df_race['Id2'], df_race[feature]))
    dict_race_year[feature + "_%02d" % (i + 6,)] = dict_race

# add each dict in dict to msa df
for key in dict_race_year:
    msa_df[key] = pd.Series(dict_race_year[key])

# cleaning column names
msa_df.columns = msa_df.columns.str.replace("Total: - ", '')
msa_df.columns = msa_df.columns.str.replace('; ', '')
msa_df = msa_df.rename(columns=lambda x: x.strip())
msa_df.columns = msa_df.columns.str.replace(r"\s+(\.*\)", "")

#
msa_df.head()

```

```

Out[19]:

```

	msa_name	pop_06	pop_07	pop_08	pop_09	\	
10180	Abilene, TX	158548	159439	160012	160266		
10380	Aguadilla-Isabela-San Sebastián, PR	336502	335201	339193	342495		
10420	Akron, OH	700943	699356	698553	699935		
10500	Albany, GA	165062	162767	163074	164238		
10580	Albany-Schenectady-Troy, NY	850957	853358	853919	857592		
	pop_10	pop_11	pop_12	pop_13	pop_14	...	\
10180	164941	165858	167800	168144	166900	...	

10380	305988	304633	304727	327847	322079	...
10420	702951	701456	702262	705686	703825	...
10500	162659	161617	155019	156277	152596	...
10580	870832	871478	874646	877905	880167	...

Native Hawaiian and Other Pacific Islander alone_Percent_15 \	
10180	0.000000
10380	0.000000
10420	0.000284
10500	0.000025
10580	0.000442

Some other race alone_Percent_15		Two or more races:_Percent_15 \	
10180	0.096624	0.033951	
10380	0.177460	0.044009	
10420	0.003466	0.029231	
10500	0.001917	0.017504	
10580	0.013795	0.032402	

White alone_Percent_16		Black or African American alone_Percent_16 \	
10180	0.781371	0.082559	
10380	0.704982	0.034152	
10420	0.819179	0.120844	
10500	0.421472	0.539284	
10580	0.827320	0.075000	

American Indian and Alaska Native alone_Percent_16 \	
10180	0.007433
10380	0.000588
10420	0.000964
10500	0.001692
10580	0.002774

Asian alone_Percent_16 \	
10180	0.020701
10380	0.000000
10420	0.028414
10500	0.009882
10580	0.043951

Native Hawaiian and Other Pacific Islander alone_Percent_16 \	
10180	0.000047
10380	0.000165
10420	0.000212
10500	0.000734
10580	0.000352

Some other race alone_Percent_16		Two or more races:_Percent_16	
----------------------------------	--	-------------------------------	--

10180	0.078637	0.029252
10380	0.226595	0.033519
10420	0.002700	0.027686
10500	0.010065	0.016871
10580	0.017621	0.032982

[5 rows x 408 columns]

Now that we have all of the Census data as one dataframe, we can split them by year. We do this below, creating a dictionary of dataframes with key = year.

```
In [20]: dict_df_year = {}

for i in range(0, 11):
    num = "_%02d" % (i + 6,)
    filter_col = [col for col in msa_df if col.endswith(num)]
    filter_col.append('msa_name')

    dict_df_year[2006 + i] = msa_df[filter_col]
```

An example:

```
In [21]: dict_df_year[2006].head()
```

```
Out[21]:
```

	pop_06	Percent_male_06	Percent_female_06	AGE - Under 5 years_06 \
10180	158548	0.497717	0.502283	7.2
10380	336502	0.495349	0.504651	6.2
10420	700943	0.481664	0.518336	5.9
10500	165062	0.476015	0.523985	7.8
10580	850957	0.485577	0.514423	5.4

	AGE - 5 to 9 years_06	AGE - 10 to 14 years_06 \
10180	6.4	7.3
10380	6.9	8.0
10420	6.3	6.9
10500	7.7	7.3
10580	5.9	6.3

	AGE - 15 to 19 years_06	AGE - 20 to 24 years_06 \
10180	8.3	8.7
10380	7.7	7.0
10420	7.2	7.2
10500	10.0	6.0
10580	7.6	7.4

	AGE - 25 to 29 years_06	AGE - 30 to 34 years_06 \
10180	6.8	6.2
10380	7.0	7.0
10420	6.3	5.9

10500	6.3	6.0
10580	6.8	6.1

	...	poverty_06	median_income_06	\
10180	...	15.8	39784.0	
10380	...	57.0	11717.0	
10420	...	12.7	44507.0	
10500	...	22.8	35515.0	
10580	...	9.8	53202.0	

	White alone_Percent_06	Black or African American alone_Percent_06	\
10180	0.741838	0.068295	
10380	0.896527	0.019961	
10420	0.844866	0.116880	
10500	0.485957	0.494136	
10580	0.867237	0.070216	

	American Indian and Alaska Native alone_Percent_06	\
10180	0.003570	
10380	0.003123	
10420	0.002129	
10500	0.001557	
10580	0.002378	

	Asian alone_Percent_06	\
10180	0.014292	
10380	0.000758	
10420	0.017715	
10500	0.006355	
10580	0.030761	

	Native Hawaiian and Other Pacific Islander alone_Percent_06	\
10180	0.000000	
10380	0.000000	
10420	0.000564	
10500	0.000000	
10580	0.000159	

	Some other race alone_Percent_06	Two or more races:_Percent_06	\
10180	0.148264	0.023740	
10380	0.058686	0.020945	
10420	0.006270	0.011577	
10500	0.003453	0.008542	
10580	0.014440	0.014809	

	msa_name
10180	Abilene, TX
10380	Aguadilla-Isabela-San Sebastián, PR

10420	Akron, OH
10500	Albany, GA
10580	Albany-Schenectady-Troy, NY

[5 rows x 38 columns]

2 Cleaning and Merging

Now we need to merge the FBI data with the Census data. We find that the FBI data uses different naming conventions from the Census Data, so we need to match each FBI df's MSA name with the Census df's MSA name, replacing the FBI names when we find a match. Below we define a few functions to help us do this. Our `match_name` function checks to see how similar a two MSA names are using the `fuzzywuzzy` package. We use a combination of scoring to return the best match.

In [22]: *# functions for splitting, cleaning, and matching names*

```
# cleaning the names
def clean_split_names(some_list):
    split_list = []
    for i in some_list:
        split_list.append(i.split(' '))
    for i in range(1, len(split_list)):
        split_list[i][0].replace(' ', '')
        split_list[i][1].replace(' ', '')

        if len(split_list[i]) > 2:
            if split_list[i][-1] == 'M.D.':

                split_list[i] = split_list[i][: -1]

            if split_list[i][-1] == '':
                split_list[i] = split_list[i][: -1]
                split_list[i][-1] = split_list[i][-1].replace(',', '')
            split_list[i][-1].replace(',', '')
    return split_list

def join_list(some_list):
    joined = []
    for i in some_list:
        joined.append(' '.join(i))
    return joined

def match_name(name, list_names, min_score = 0):
```

```

# -1 score incase we don't get any matches
max_score = -1

# Returning empty name for no match as well
max_name = ""

# Iterating over all names in the other
for name2 in list_names:

    score3 = fuzz.ratio(''.join(name), ''.join(name2))

    # make sure they are from the same state
    if name2[-1].replace(',','') == name[-1].replace(',',''):

        # get similiarity scores
        score1 = fuzz.ratio(''.join(name), ''.join(name2))
        score2 = fuzz.partial_ratio(''.join(name), ''.join(name2))

        # and take the average of them
        score_av = (score1 + score2)/ 2

        if score3 > score_av:
            score = score3
        if score3 <= score_av:
            score = score_av

        # Checking if we are above our threshold and have a better score
        if (score > min_score) & (score > max_score):
            max_name = name2
            max_score = score
return (''.join(max_name), max_score)

```

In [23]: # cleaning names

```

for i in range(2006,2017):
    murder_stats[i]['MSA'] = murder_stats[i]['MSA'].str.replace(' M.S.A.', '')
    murder_stats[i]['MSA'] = murder_stats[i]['MSA'].str.replace('\d+', '')
    murder_stats[i]['MSA'] = murder_stats[i]['MSA'].str.replace('\ue83a', ' ')
    murder_stats[i]['MSA'] = murder_stats[i]['MSA'].str.replace('\nM.S.A.', ' ')
    murder_stats[i]['MSA'] = murder_stats[i]['MSA'].str.replace('ãĚ.', ' ')
    murder_stats[i]['MSA'] = murder_stats[i]['MSA'].str.replace('urfrees', ' ')
    murder_stats[i]['MSA'] = murder_stats[i]['MSA'].str.replace('M.S.A.', '')
    murder_stats[i]['MSA'] = murder_stats[i]['MSA'].str.replace('M.D.', '')

murder_stats[i].rename(columns={'Murder': 'Murder ' + str(i)}, inplace = True)

```



```

msa_year = murder_stats[i]['MSA'].tolist()
clean_names = clean_split_names(msa_year)
join_names = join_list(clean_names)
murder_stats[i]['MSA'] = join_names

```

Below we match each FBI dataframe from a given year with the Census dataframe from the same year by first merging the dataframe, and finding the FBI MSA names that did not match and the Census MSA names that were not matched to. We use these names to find further matches, replacing the FBI names with the Census names. Last after we have replaced names, we merge the dataframe together again, and drop all the columns that do not match (some MSA's only appear in one dataset). We print the dimensions of the merged dfs before matching and after matching to show how well our matching works.

```

In [24]: # matching the names together and merging the dataframes
for i in range(2006, 2017):

    # merge without any matching
    data = dict_df_year[i].merge(murder_stats[i],
                                left_on='msa_name',
                                right_on='MSA',
                                how = 'outer',
                                #since they are the same name, we need a suffix
                                )

    # look at the shape (we want to see if this will decrease with matching)
    print(data.shape)

    # find the census names that haven't been matched
    census = data[data['MSA'].isnull()]['msa_name'].tolist()
    census_split = clean_split_names(census)

    # find the fbi names that haven't been matched
    fbi = data[data['msa_name'].isnull()]['MSA'].tolist()
    fbi_split = clean_split_names(fbi)

    dict_list = []
    names_replace = []

    # for every unmatched fbi names
    for name in fbi_split:

        # return the best matched string in the census unmatched list if the match
        match = match_name(name, census_split, 66.7)

        # put into a dict
        dict_ = {}
        dict_.update({"FBI" : ' '.join(name)})
        dict_.update({'Census' : match[0]})
        dict_.update({"Score" : match[1]})

```

```

dict_list.append(dict_)

names_replace.append(match[0])

# put into a df
merge_table = pd.DataFrame(dict_list)

# isolate the matches since unmatched will have a score of -1 by default
replace_table = merge_table[merge_table['Score'] > 5][['FBI', 'Census']]

# put matches in a dict with key: the str of fbi value: its census name match
dict_replace = dict(zip(replace_table.FBI, replace_table.Census))

# replace its name with the Census name
murder_stats[i]['MSA'].update(murder_stats[i]['MSA'].map(dict_replace))

# merge again after matching
data2 = dict_df_year[i].merge(murder_stats[i],
                              left_on = 'msa_name',
                              right_on = 'MSA',
                              how = 'inner',
                              )
data2 = data2[pd.notnull(data2['pop_' + '%02d' % (i - 2000)])]

# this number should be now lower
print(data2.shape, 'Should be less')

# update the dictionary of dataframes
dict_df_year[i] = data2

```

```

(445, 40)
(324, 40) Should be less
(456, 40)
(327, 40) Should be less
(455, 40)
(323, 40) Should be less
(471, 40)
(344, 40) Should be less
(461, 40)
(342, 40) Should be less
(472, 40)
(339, 40) Should be less
(508, 40)
(318, 40) Should be less
(510, 40)
(335, 40) Should be less
(502, 40)

```

```
(335, 40) Should be less
(497, 40)
(334, 40) Should be less
(507, 40)
(337, 40) Should be less
```

An example of the merged dataframe:
We drop the additional name column.

```
In [25]: for i in range(2006, 2017):
         dict_df_year[i] = dict_df_year[i].drop('MSA', axis = 1)
```

We now make the id's a column and reindex:

```
In [26]: # added ID's and saved the file
         for i in range(2006, 2017):
             unique_id = pd.DataFrame(np.c_[msa_df.index, msa_df['msa_name']], columns = ['ID', 'MSA'])
             dict_df_year[i] = dict_df_year[i].merge(unique_id, left_on = 'msa_name', right_on = 'MSA')
```

2.1 BEA Data For Feature GDP per Capital

We add one more feature, gdp per capital for each year. This data was taken from the BEA.

```
In [27]: df_gdp = pd.read_csv('gdp_per_capita.csv')

         df_gdp.columns = ['IDs', 'MSA', 'gdp_06', 'gdp_07', 'gdp_08', 'gdp_09', 'gdp_10', 'gdp_11', 'gdp_12', 'gdp_13', 'gdp_14', 'gdp_15', 'gdp_16']
         df_gdp = df_gdp.drop(0)
         df_gdp.head()
```

```
Out[27]:
```

	IDs	MSA	gdp_06	gdp_07	\
1	10180	Abilene, TX (Metropolitan Statistical Area)	33978	34883	
2	10420	Akron, OH (Metropolitan Statistical Area)	42081	42272	
3	10500	Albany, GA (Metropolitan Statistical Area)	32657	31885	
4	10540	Albany, OR (Metropolitan Statistical Area)	31331	29552	
5	10580	Albany-Schenectady-Troy, NY (Metropolitan Statistical Area)	49549	48662	

	gdp_08	gdp_09	gdp_10	gdp_11	gdp_12	gdp_13	gdp_14	gdp_15	gdp_16
1	35663	33920	34004	33964	35406	37550	39776	39631	38385
2	42633	40667	41138	40777	40687	41654	43750	46022	48195
3	31376	31848	31036	30332	30578	30902	30005	29235	29073
4	28153	27756	27049	27368	28171	27930	27520	28111	29328
5	48504	49279	49716	49028	49548	49823	50149	51328	51755

```
In [28]: # add feature to each dictionary
         for i in range(2006, 2017):
```

```
dict_df_year[i] = dict_df_year[i].merge(df_gdp[['IDs', 'gdp_' + '%02d' % (i - 2000)],
                                         left_on = 'ID',
                                         right_on = 'IDs',
                                         how = 'inner',
                                         )
dict_df_year[i] = dict_df_year[i].drop('IDs', axis = 1)
```

```
In [29]: # add to dict msa_df for plotting later
dict_gdp_year = {}

for i in range(2006, 2017):
    dict_gdp_year['gdp_' + '%02d' % (i - 2000)] = dict(zip(df_gdp['IDs'], df_gdp['gdp_']

for key in dict_gdp_year:
    msa_df[key] = pd.Series(dict_gdp_year[key])
```

3 EDA

We make two columns for EDA. The first is `murder_category` which categorizes each year's murder counts by quartile. Up to the 25th percentile of 2016 numbers, murders are categorized with a 0 for low murder, from 25th to 75th they are categorized with a 1 for medium murder, and from 75th quartile and over, murders are categorized with a 2 for high.

We make another column called `'murder_rate'`, which takes the murder counts of each year and divides it by the population. Then we multiply it by 10,000 to make it murder per 10,000 people

```
In [30]: quantile = dict_df_year[2016]['Murder 2016'].quantile([0.25,0.5,0.75])
print(quantile)

how_murder = lambda x:((0, 1)[x > quantile[.25]], 2)[x > quantile[.75]]

for i in range(2006, 2017):
    dict_df_year[i]['murder_category'] = dict_df_year[i]['Murder ' + str(i)].apply(how_
dict_df_year[i]['pop_' + '%02d' % (i - 2000)] = dict_df_year[i]['pop_' + '%02d' % (
dict_df_year[i]['murder_rate'] = dict_df_year[i]['Murder ' + str(i)] / dict_df_year
dict_df_year[i]['murder_rate'] = dict_df_year[i]['murder_rate'] * 100000

0.25    4.0
0.50    9.0
0.75   26.0
Name: Murder 2016, dtype: float64
```

3.1 Feature Trend Lines

In order to determine how the features have changed over time, we select four MSA's. In particular, Chicago-Naperville-Joliet, IL-IN-WI, Rochester, Duluth, MN-WI, Philadelphia-Camden-

Wilmington, PA-NJ-DE-MD, Panama City-Lynn Haven, FL from 2016 counts. One of the MSA's is in the low murder category, another is the medium murder category, and the other two are in the high murder category.

In [31]: *### pick some cities of interest based on murder rates/row choice in 2016*

```
def find_nearest(array, value):
    idx = (np.abs(array - value)).argmin()
    return idx

murders_16 = dict_df_year[2016]['Murder 2016']

maxidx = np.argmax(murders_16)
minidx = np.argmin(murders_16)

mediumidx = find_nearest(murders_16, (max(murders_16) + min(murders_16)) / 2)
sub_16 = dict_df_year[2016].iloc[[maxidx, 86, mediumidx, 211], :]
sub_ids = sub_16['ID'].tolist()

sub_16
```

```
Out[31]:
```

	pop_16	Percent_male_16	Percent_female_16	AGE - Under 5 years_16 \
54	9512968	0.489620	0.510380	6.1
86	279227	0.505678	0.494322	5.1
214	6070500	0.483860	0.516140	5.9
211	200893	0.504821	0.495179	5.9

	AGE - 5 to 9 years_16	AGE - 10 to 14 years_16	AGE - 15 to 19 years_16 \
54	6.4	6.7	6.7
86	5.1	6.2	6.9
214	6.1	6.2	6.5
211	6.7	4.8	6.2

	AGE - 20 to 24 years_16	AGE - 25 to 29 years_16 \
54	6.8	7.2
86	8.1	5.8
214	6.6	7.4
211	6.3	7.1

	AGE - 30 to 34 years_16	...	Asian alone_Percent_16 \
54	7.1	...	0.064895
86	5.8	...	0.010687
214	6.8	...	0.058514
211	6.5	...	0.027920

	Native Hawaiian and Other Pacific Islander alone_Percent_16 \
54	0.000387
86	0.000161

214		0.000357
211		0.000169

	Some other race alone_Percent_16	Two or more races:_Percent_16	\
54	0.087865	0.026341	
86	0.004136	0.023132	
214	0.034570	0.028379	
211	0.005306	0.030763	

	msa_name	Murder 2016	ID	gdp_16	\
54	Chicago-Naperville-Joliet, IL-IN-WI	958.0	16980	59810	
86	Duluth, MN-WI	3.0	20260	39814	
214	Philadelphia-Camden-Wilmington, PA-NJ-DE-MD	480.0	37980	62817	
211	Panama City-Lynn Haven, FL	9.0	37460	36374	

	murder_category	murder_rate
54	2	10.070464
86	0	1.074395
214	2	7.907092
211	1	4.479997

[4 rows x 43 columns]

We define a function that will plot the feature name over 2006-2016 for specified MSAs


```
In [32]: # use msa just for the feature trends
# function can take ids of MSA's and features and plot the trend of that feature in time
def trend_plot(feature_name, sub_ids, title):
    pc_temp = []
    for i in range(2006, 2017):
        pc_temp.append(msa_df.loc[sub_ids].columns.get_loc(feature_name + '%02d' % (i - 2006)))

    df_temp = msa_df.loc[sub_ids].iloc[:, pc_temp]
    df_temp.loc[len(df_temp)] = np.arange(2006, 2017, 1)
    df_temp = df_temp.T
    new_cols = msa_df.loc[sub_ids]['msa_name'].tolist()
    new_cols.append('year')
    df_temp.columns = new_cols
    plt.figure(figsize = (12, 7))


    for i in range(len(sub_ids)):
        plt.plot(df_temp['year'].values, df_temp.iloc[:, i].values, label = df_temp.columns[i])
    plt.legend(loc = 'best')
    plt.xlabel('Years')
    plt.ylabel(feature_name.replace("_", " "))
    plt.title(title)
```

We plot some of the features below:


```
In [33]: trend_plot('Divorced_', sub_ids, 'Percent Divorced Over Time')
```



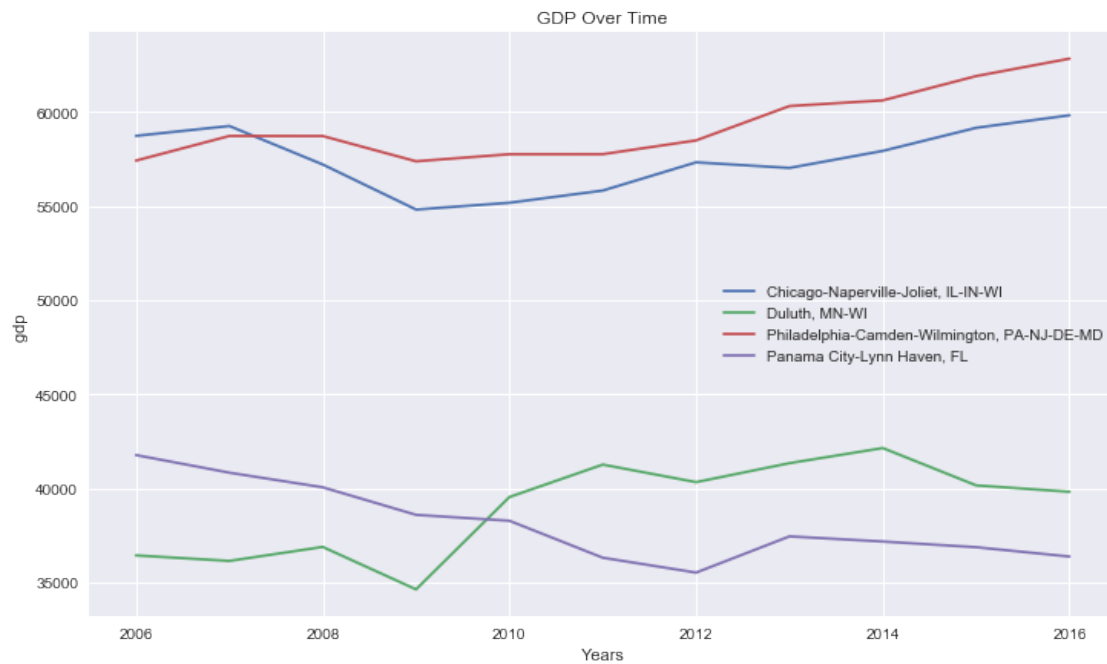
```
In [34]: trend_plot('poverty_', sub_ids, 'Percent Poverty Over Time')
```



```
In [35]: trend_plot('Black or African American alone_Percent_', sub_ids, 'Percent Black Over Time')
```



```
In [36]: trend_plot('gdp_', sub_ids, 'GDP Over Time')
```



3.2 Murder Over Time

```
In [37]: sub_years = {}
        for i in range(2006,2017):
            sub_years[i] = dict_df_year[i].loc[dict_df_year[i]['ID'].isin(sub_ids)]
```

```
In [38]: sub_years[2008].head()
```

```
Out[38]:
```

	pop_08	Percent_male_08	Percent_female_08	AGE - Under 5 years_08	\
76	275336	0.494599	0.505401	5.5	
208	163946	0.490924	0.509076	6.6	
210	5838471	0.483937	0.516063	6.5	

	AGE - 5 to 9 years_08	AGE - 10 to 14 years_08	AGE - 15 to 19 years_08	\
76	5.7	5.3	7.7	
208	6.9	6.1	6.2	
210	6.4	6.6	7.5	

	AGE - 20 to 24 years_08	AGE - 25 to 29 years_08	\
76	8.4	6.6	
208	5.4	6.8	
210	6.7	5.9	

	AGE - 30 to 34 years_08	...	Asian alone_Percent_08	\
76	5.4	...	0.006505	
208	6.4	...	0.017579	
210	5.9	...	0.043522	

	Native Hawaiian and Other Pacific Islander alone_Percent_08	\
76	0.000000	
208	0.002049	
210	0.000195	

	Some other race alone_Percent_08	Two or more races:_Percent_08	\
76	0.000701	0.017161	
208	0.004264	0.040751	
210	0.028367	0.017605	

	msa_name	Murder 2008	ID	gdp_08	\
76	Duluth, MN-WI	4.0	20260	36885	
208	Panama City-Lynn Haven, FL	11.0	37460	40052	
210	Philadelphia-Camden-Wilmington, PA-NJ-DE-MD	530.0	37980	58708	

	murder_category	murder_rate
76	0	1.452770
208	1	6.709526
210	2	9.077719

[3 rows x 43 columns]


```

In [39]: # looks at sub years and checks whether all ids are there and extracts murders
# if some id's are missing, murders for those are set to zero (e.g. Chicago from 2006 t
def find_yw_data(sub_years, sub_ids):
    main = []
    for i in range(2006,2017):
        temp = []
        for j in sub_ids:
            if sum((sub_years[i]['ID'])== j) == 1:
                temp.append(sub_years[i][sub_years[i]['ID'] == j]['Murder ' + str(i)].va
            else:
                temp.append(0)
        main.append(temp)
    return(main)

In [40]: # use msa just for the feature trends
# plots murder trends given MSA indices
def murder_trend_plot(sub_ids):
    murder_years = np.array(find_yw_data(sub_years,sub_ids))
    new_cols = msa_df.loc[sub_ids]['msa_name'].tolist()
    x = np.arange(2006,2017,1)
    plt.figure(figsize=(16, 8))
    for i in range(len(sub_ids)):
        plt.plot(x,murder_years[:,i],label = new_cols[i])
    plt.legend(loc='best')
    plt.xlabel('Years')
    plt.ylabel('Murders')
    plt.title('Murder Counts Over Time')

In [41]: murder_trend_plot(sub_ids)

```



3.3 Murder by Feature


```
In [42]: # use msa just for the feature trends
        # scatter plot of feature against the murder with data from all years

def feature_murder_plot(feature_name, sub_ids, title):
    pc_temp = []
    sub_years = {}
    for i in range(2006, 2017):
        sub_years[i] = dict_df_year[i].loc[dict_df_year[i]['ID'].isin(sub_ids)]


    for i in range(2006, 2017):
        pc_temp.append(msa_df.loc[sub_ids].columns.get_loc(feature_name + '%02d' % (i -
murder_years = np.array(find_yw_data(sub_years, sub_ids))
        # index all rows with interesting sub ids and feature of interest columns
        df_temp = msa_df.loc[sub_ids].iloc[:, pc_temp]
        # add row for years
        df_temp.loc[len(df_temp)] = np.arange(2006, 2017, 1)
        df_temp = df_temp.T
        new_cols = msa_df.loc[sub_ids]['msa_name'].tolist()
        new_cols.append('year')
        df_temp.columns = new_cols
        plt.figure(figsize=(12, 6))
        for i in range(len(sub_ids)):
            plt.scatter(df_temp.iloc[:, i].values, murder_years[:, i], label = new_cols[i])
        plt.legend(loc='best')
        plt.ylabel('# of Murders')
        plt.xlabel(feature_name.replace("_", " "))
        plt.title(title)
```

We plot some of the features below:


```
In [43]: feature_murder_plot('Percent_male_', sub_ids, 'Murders by Percentage of Males')
```



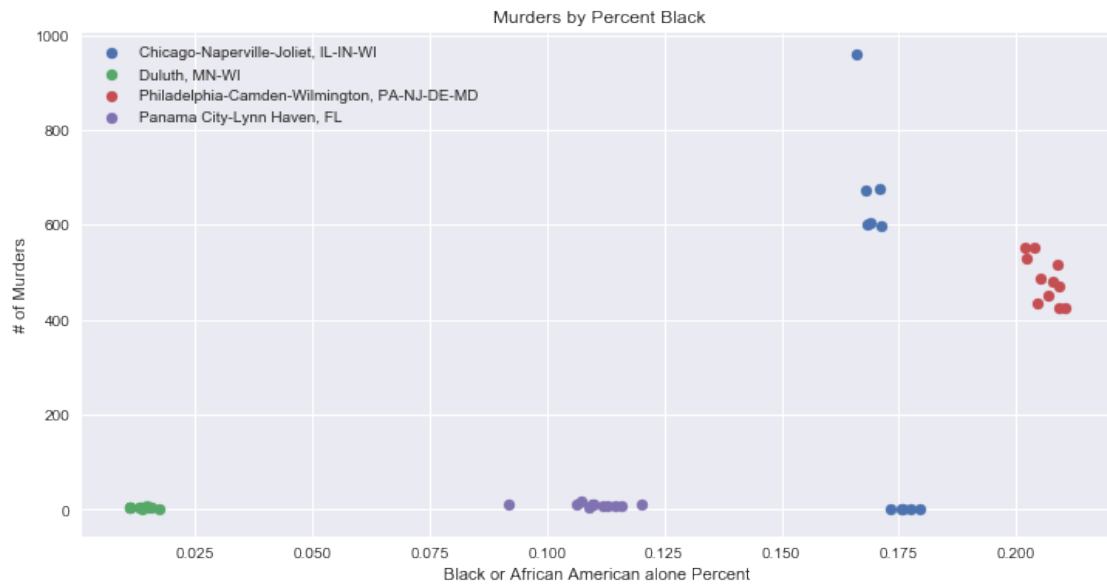
```
In [44]: feature_murder_plot('Divorced_', sub_ids, 'Murders by Percentage Divorced')
```



```
In [45]: feature_murder_plot('poverty_', sub_ids, 'Murders by Poverty')
```



In [46]: `feature_murder_plot('Black or African American alone_Percent_', sub_ids, 'Murders by Pe`



3.4 Murder on Map of US

We use Geopy to find the longitude and latitude of each MSA and plot them on a map'

```

In [47]: ## find lat and long for each msa
geolocator = Nominatim()
# lat_long = []
# for row in unique_id['msa_name']:
#     if geolocator.geocode(row) == None:
#         lat_long.append(np.nan)
#     if geolocator.geocode(row) != None:
#         lat_long.append(geolocator.geocode(row)[1])

# # save as file
# with open('outfile', 'wb') as fp:
#     pickle.dump(lat_long, fp)

# read in file
with open('outfile', 'rb') as fp:
    llng = pickle.load(fp)

# split by latitude and longitude
lat = []
long = []
for i in llng:
    if type(i) is float:
        lat.append(np.nan)
        long.append(np.nan)
    if type(i) is tuple:
        lat.append(i[0])
        long.append(i[1])

lat_long_df = pd.DataFrame([unique_id['ID'].tolist(), lat, long]).T

# some of the lat_longs were wrong
lat_long_df.columns = ['ID', 'latitude', 'longitude']

In [48]: # add to each df
position_murder_df = {}
for i in range(2006, 2017):
    position_murder_df[i] = dict_df_year[i].merge(lat_long_df, left_on = 'ID', right_on=

In [49]: # some lat and long were wrong so we drop them
position_murder_df[2006] = position_murder_df[2006][position_murder_df[2006]['msa_name']
position_murder_df[2016] = position_murder_df[2016][position_murder_df[2016]['msa_name']

In [50]: plotly.tools.set_credentials_file(username='jeanettejin', api_key='65684eh97DVINyUzjc

data = [dict(
    lat = position_murder_df[2006]['latitude'],
    lon = position_murder_df[2006]['longitude'],
    text = position_murder_df[2006]['Murder 2006'].astype(str) + ' count',

```

```

marker = dict(
    size = 10,
    color = position_murder_df[2006]['Murder 2006'],
    colorsrc = "jeanettejin:2:236ccc",
    opacity = 1.0,
    colorscale = [[0, 'rgb(49,54,149)'], [1./10000, 'rgb(69,117,180)'], [1./1000,
    colorbar = dict(
        thickness = 10,
        titleside = "right",
        outlinecolor = "rgb(68, 68, 68, 0)",
        ticklen = 3,
        showticksuffix = 'first',
        ticksuffix = " count",
    ),
),
type = 'scattergeo',
) ]

layout = dict(
    geo = dict(
        scope = 'north america',
        showland = True,
        landcolor = "rgb(212, 212, 212)",
        subunitcolor = "rgb(255, 255, 255)",
        countrycolor = "rgb(255, 255, 255)",
        showlakes = True,
        lakecolor = "rgb(255, 255, 255)",
        showsubunits = True,
        showcountries = True,
        resolution = 50,
        projection = dict(
            type = 'conic conformal',
            rotation = dict(
                lon = -100
            )
        ),
        lonaxis = dict(
            showgrid = True,
            gridwidth = 0.5,
            range= [ -140.0, -55.0 ],
            dtick = 5
        ),
        lataxis = dict (
            showgrid = True,
            gridwidth = 0.5,
            range= [ 20.0, 60.0 ],
            dtick = 5

```

```

    ),
    title = 'Murder Counts in US 2006',
)
fig = {'data': data, 'layout': layout }
py.iplot(fig, filename = 'Murder_2006')

```

Out[50]: <plotly.tools.PlotlyDisplay object>

In [51]: `plotly.tools.set_credentials_file(username='jeanettejin', api_key='65684eh97DVINyUzjo')`

```

data = [dict(
    lat = position_murder_df[2016]['latitude'],
    lon = position_murder_df[2016]['longitude'],
    text = position_murder_df[2016]['Murder 2016'].astype(str) + ' count',

    marker = dict(
        size = position_murder_df[2016]['Murder 2016'] / 10,
        color = 'red',
        colorsrc = "jeanettejin:2:236ccc",
        opacity = 1.0,
    ),
    type = 'scattergeo',
) ]

layout = dict(
    geo = dict(
        scope = 'north america',
        showland = True,
        landcolor = "rgb(212, 212, 212)",
        subunitcolor = "rgb(255, 255, 255)",
        countrycolor = "rgb(255, 255, 255)",
        showlakes = True,
        lakecolor = "rgb(255, 255, 255)",
        showsubunits = True,
        showcountries = True,
        resolution = 50,
        projection = dict(
            type = 'conic conformal',
            rotation = dict(
                lon = -100
            )
        ),
        lonaxis = dict(
            showgrid = True,
            gridwidth = 0.5,
            range= [ -140.0, -55.0 ],
            dtick = 5

```



```

    ),
    lataxis = dict (
        showgrid = True,
        gridwidth = 0.5,
        range= [ 20.0, 60.0 ],
        dtick = 5
    )
),
title = 'Murder Counts in US 2016',
)
fig = {'data': data, 'layout': layout }
py.iplot(fig, filename = 'Murder_2016')

```

Out [51]: <plotly.tools.PlotlyDisplay object>

3.5 Histograms of Murders Across Years

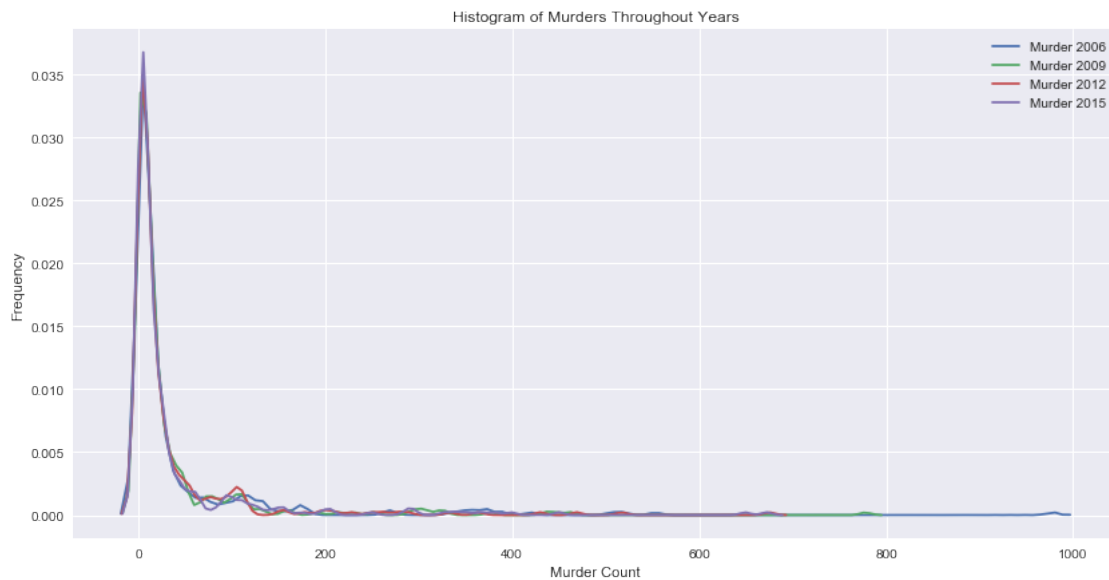
We can see that most MSA have pretty low murder counts

```

In [52]: fig, ax = plt.subplots(figsize=(14, 7))
        ax.set(xlabel = 'Murder Count', ylabel = 'Frequency')
        ax.set_title("Histogram of Murders Throughout Years")

        for i in range(2006, 2017, 3):
            sns.kdeplot(dict_df_year[i]['Murder ' + str(i)], ax = ax)

```



```

In [53]: fig, ax = plt.subplots(figsize=(14, 7))
        ax.set(xlabel = 'Murder Count', ylabel = 'Frequency')

```

```
ax.set_title("Histogram of Murders Throughout Years")
```

```
for i in range(2006, 2017, 3):
    sns.kdeplot(dict_df_year[i]['murder_rate'], ax = ax)
```

```
plt.legend(['Murder per 10,000 2006', 'Murder per 10,000 2009', 'Murder per 10,000 2012'])
```

Out[53]: <matplotlib.legend.Legend at 0x1242037b8>



3.6 Feature vs Murder or Murder / Population by Group Through Time

We define a function that plots murder's by feature in certain circumstances, plotted throughout various years. If option is 0 we plot murder by feature for various years. If option is 1, we plot murder by feature and differentiate each point by weather it's murder_category is low, medium, or high across multiple years. If option is 2, we plot murder divided by population by feature across multiple years.

```
In [54]: def feature_x_time(title, feature, option = 0):

    if option is 0:
        fig, ax = plt.subplots(1, 4, figsize=(20, 10))
        plt.subplots_adjust(top = .94)
        fig.suptitle(title)
        for i in range(2006, 2017, 3):
            ind = int((i - 2006) / 3)
```

```

sns.regplot(feature + '%02d' % (i - 2000), 'Murder ' + str(i), dict_df_year
sns.set(font_scale = .7)

if option is 1:
    for i in range(2006, 2017, 3):
        ind = int((i - 2006) / 3)
        sns.lmplot(feature + '%02d' % (i - 2000), 'Murder ' + str(i), hue = "murder


if option is 2:
    fig, ax = plt.subplots(1, 4, figsize=(20, 10))
    plt.subplots_adjust(top = .94)
    fig.suptitle(title)

    for i in range(2006, 2017, 3):
        ind = int((i - 2006) / 3)
        sns.regplot(feature + '%02d' % (i - 2000), 'murder_rate', dict_df_year[i],
ax[ind].set(ylim=(0, 6))

```

We can see that murder and population seem to have a strong linear relationship. For the remainder of the variables, we also include option 2, so we can see how each feature varies with murder/population


```
In [55]: feature_x_time("Murder by Population", 'pop_')
```





```
In [56]: feature_x_time("Percent Divorced and Murder Across Time", "Divorced_")
```




```
In [57]: feature_x_time("Percent Divorced and Murder Across Time", "Divorced_", option = 1)
```











In [58]: feature_x_time("Percent Divorced and Murder per 100,000 Across Time", "Divorced_", opti





```
In [59]: feature_x_time('Poverty and Murder Across Time', 'poverty_')
```




```
In [60]: feature_x_time('Poverty and Murder Across Time', 'poverty_', option = 1)
```






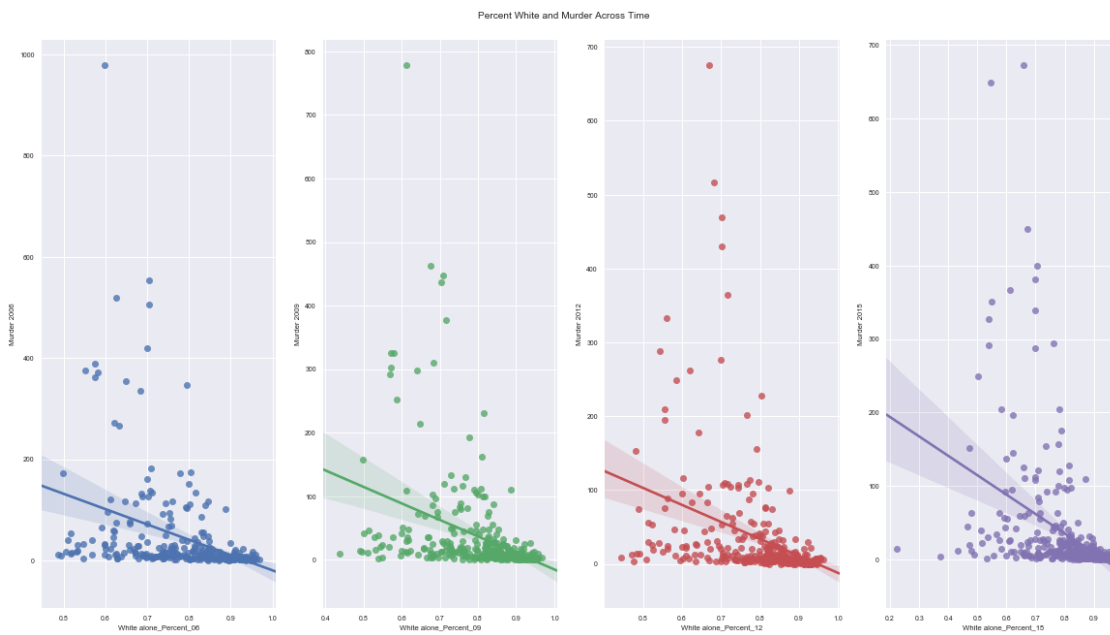





```
In [61]: feature_x_time('Poverty and Murder per 100,000 Across Time', 'poverty_', option = 2)
```





```
In [62]: feature_x_time('Percent White and Murder Across Time', 'White alone_Percent_')
```




```
In [63]: feature_x_time('Percent White and Murder Across Time', 'White alone_Percent_', option =
```










```
In [64]: feature_x_time('Percent White and Murder per 100,000 Across Time', 'White alone_Percent
```




```
In [65]: feature_x_time("GDP and Murder Across Time", 'gdp_')
```




```
In [66]: feature_x_time("GDP and Murder Across Time", 'gdp_', option = 1)
```










```
In [67]: feature_x_time("GDP and Murder per 100,000 Across Time", 'gdp_', option = 2)
```




3.7 Feature-Feature Correlation

Looks like multicollinearity will be an issue. We will address this through feature selection

```
In [68]: fig, ax = plt.subplots(figsize=(10, 8))
```

```
# corr heat map
sns.set(font_scale = 1.5)
sns.heatmap(dict_df_year[2006].corr(), ax = ax)
```

```
Out[68]: <matplotlib.axes._subplots.AxesSubplot at 0x11f506860>
```

4 Modeling

Before we do our model, let's first collapse some of our age columns.

```
In [69]: # collapse age columns and drop specific columns
for i in range(2006, 2017):
    dict_df_year[i]['under_18_' + '%02d' % (i - 2000)] = dict_df_year[i]['AGE - Under 5 years_' + '%02d' % (i - 2000)]
    dict_df_year[i]['20_to_40_' + '%02d' % (i - 2000)] = dict_df_year[i]['AGE - 20 to 44 years_' + '%02d' % (i - 2000)]
    dict_df_year[i]['40_to_60_' + '%02d' % (i - 2000)] = dict_df_year[i]['AGE - 45 to 64 years_' + '%02d' % (i - 2000)]
    dict_df_year[i]['60_above_' + '%02d' % (i - 2000)] = dict_df_year[i]['AGE - 65 to 84 years_' + '%02d' % (i - 2000)]
    dict_df_year[i] = dict_df_year[i].drop(['AGE - Under 5 years_' + '%02d' % (i - 2000),
                                           'AGE - 20 to 44 years_' + '%02d' % (i - 2000),
                                           'AGE - 45 to 64 years_' + '%02d' % (i - 2000),
                                           'AGE - 65 to 84 years_' + '%02d' % (i - 2000)])
```

```
In [70]: dict_avg = {}
```

```
for i in range(2006, 2017):
    dict_avg[i] = pd.DataFrame.copy(dict_df_year[i])
```

```

In [71]: # function that standarizes df with 2006 min, max and range and returns the x_train and
def preprocess(year, columns, y_name):

    # empty lists for column names
    columns_year = []
    columns_06 = []
    year_2 = '%02d' % (year - 2000)
    year_06 = '%02d' % 6

    # make columns names
    for column in columns:
        columns_year.append(column + year_2)
        columns_06.append(column + year_06)

    # df
    df = dict_df_year[year].dropna()

    # split
    x_train = df[columns_year]
    y_train = df[[y_name]]

    # scale df
    df_scale = dict_df_year[2006].dropna()

    # select columns
    columns_scale = df_scale[columns_06]
    x_columns = x_train.columns.tolist()

    # standardize
    scaler = StandardScaler().fit(columns_scale)
    x_train[x_columns] = scaler.transform(x_train)

    return x_train, y_train

# function that adds polynomial terms to x_train and x_test
# polfeatures
def more_terms(x_train, x_test, degree):
    columns = base_col
    poly = PolynomialFeatures(degree, interaction_only = False, include_bias = False)

    x_trainpoly = poly.fit_transform(x_train)
    x_testpoly = poly.fit_transform(x_test)
    feature = poly.get_feature_names(columns)
    x_train_poly = pd.DataFrame(x_trainpoly, columns = feature)
    x_test_poly = pd.DataFrame(x_testpoly, columns = feature)

    return x_train_poly, x_test_poly

```

We make a dict for all the features in x_year and all the murder counts over population in y_year

```
In [72]: # initialize dicts
        x_year = {}
        y_year = {}

        # features
        columns = ['Percent_male_' , 'under_18_', '20_to_40_', '40_to_60_', '60_above_', 'Now m

        # for each year scale and split and put into dict
        for year in range(2006, 2017):

            xyear, yyear = preprocess(year, columns, 'murder_rate')

            x_year[year] = xyear
            y_year[year] = yyear
```

4.1 Base Model (Ridge and Lasso)

We use the gender, age, and income as our basemodel. To account for time series, for each feature and year, we obtain a coefficient and use these to fit a model to each feature. We predict the coefficient for each feature, which gives us our model for 2016. We use RidgeCV and LassoCV to model 2016 data and make predictions on the murder/population.

```
In [73]: # col
        base_col = ['Percent_male_' , 'under_18_', '20_to_40_', '40_to_60_', '60_above_', 'medi

        # initialize dicts
        x_yearb = {}
        y_yearb = {}

        # split select feature for every year and fil dict
        for year in range(2006, 2017):

            xbyear, ybyear = preprocess(year, base_col, 'murder_rate')

            x_yearb[year] = xbyear
            y_yearb[year] = ybyear
```

```
In [74]: # initialize df
        coefficients_lassob = pd.DataFrame()
        coefficients_ridgeb = pd.DataFrame()

        # constants
        constant_lassob = []
        constant_ridgeb = []
```

```

# add col names
coefficients_lassob['columns'] = base_col
coefficients_ridgeb['columns'] = base_col

# get betas for every year using lasso cv and ridge cv
for i in range(2006, 2016):
    lassob = LassoCV().fit(x_yearb[i], y_yearb[i].values.reshape(-1, 1))
    coefficients_lassob[i] = lassob.coef_

    ridgeb = RidgeCV().fit(x_yearb[i], y_yearb[i].values.reshape(-1, 1))
    coefficients_ridgeb[i] = ridgeb.coef_[0]

# constants
constant_lassob.append(lassob.intercept_)
constant_ridgeb.append(ridgeb.intercept_)

```

```

In [75]: # transpose and rename
coefficients_lassob = rename(coefficients_lassob.T)
coefficients_ridgeb = rename(coefficients_ridgeb.T)

```

```

In [76]: # show
coefficients_lassob.head()

```

```

Out[76]:      Percent_male_  under_18_    20_to_40_  40_to_60_  60_above_  median_income_
2006      -0.640029   0.882617    -0.108982   0.444397        -0      -0.0501331
2007      -0.434509   1.01228    -0.0344746   0.643091        -0      -0.379183
2008      -0.661512   0.87094        -0        0.24985        -0      -0.0917136
2009      -0.519636   0.733181  -0.000297521   0.354195        -0      -0.465388
2010      -0.654053   0.317147        -0        0.436378  -0.394344      -0.404777

```

```

In [77]: # show
coefficients_ridgeb.head()

```

```

Out[77]:      Percent_male_  under_18_    20_to_40_  40_to_60_  60_above_  median_income_
2006      -0.635164   0.697911    -0.35722   0.346686  -0.313517      -0.124067
2007      -0.434678   0.643388    -0.400622   0.363601  -0.457374      -0.394092
2008      -0.665104   0.663585    -0.309884   0.142336  -0.397116      -0.180151
2009      -0.530811   0.524614    -0.263348   0.229722  -0.345614      -0.514892
2010      -0.663024   0.344533    -0.0204428  0.487956  -0.437737      -0.444311

```

4.1.1 Base Model Visualize

LassoCV

```

In [78]: plt.scatter(np.arange(2006, 2016, 1), constant_lassob)

```

```

Out[78]: <matplotlib.collections.PathCollection at 0x11ffee518>


```



```
In [79]: sns.lmplot(x = 'index', y = 'Percent_male_', data = coefficients_lassob.reset_index(),
plt.title("Lasso: Coefficients of Percent Male over Time")
plt.ylabel('Percent Male Coefficient', fontsize=16)
plt.xlabel('Year', fontsize=16)
plt.xticks([2007, 2009, 2011, 2013, 2015])
```


```
Out[79]: ([<matplotlib.axis.XTick at 0x1242be320>,
<matplotlib.axis.XTick at 0x11f62e780>,
<matplotlib.axis.XTick at 0x11f665780>,
<matplotlib.axis.XTick at 0x1248bb3c8>,
<matplotlib.axis.XTick at 0x123ee09b0>],
<a list of 5 Text xticklabel objects>)
```

Lasso: Coefficients of Percent Male over Time



```
In [80]: sns.lmplot(x = 'index', y = 'median_income_', data= coefficients_lassob.reset_index(),
```

```
Out[80]: <seaborn.axisgrid.FacetGrid at 0x11d387b70>
```



```
In [81]: sns.lmplot(x = 'index', y = 'under_18_', data= coefficients_lassob.reset_index(), fit_r
```


```
Out[81]: <seaborn.axisgrid.FacetGrid at 0x125253c88>
```



Ridge CV


```
In [82]: plt.scatter(np.arange(2006, 2016, 1), constant_ridgeb)
```

```
Out[82]: <matplotlib.collections.PathCollection at 0x11ffcb470>
```


```
In [83]: sns.lmplot(x = 'index', y = 'Percent_male_', data = coefficients_ridgeb.reset_index(),
plt.title("Ridge: Coefficients of Percent Male over Time")
plt.ylabel('Percent Male Coefficient', fontsize=16)
plt.xlabel('Year', fontsize=16)
plt.xticks([2007, 2009, 2011, 2013, 2015]))
```

```
Out[83]: ([<matplotlib.axis.XTick at 0x11e7b20b8>,
<matplotlib.axis.XTick at 0x120d5efd0>,
<matplotlib.axis.XTick at 0x123f44278>,
<matplotlib.axis.XTick at 0x123f44f60>,
<matplotlib.axis.XTick at 0x123f3ba58>],
<a list of 5 Text xticklabel objects>)
```



```
In [84]: sns.lmplot(x = 'index', y = 'Percent_male_', data= coefficients_ridgeb.reset_index(), f
```

```
Out[84]: <seaborn.axisgrid.FacetGrid at 0x1246a8ba8>
```



```
In [85]: sns.lmplot(x = 'index', y = 'Percent_male_', data= coefficients_ridgeb.reset_index(), f
```

```
Out[85]: <seaborn.axisgrid.FacetGrid at 0x11e97a3c8>
```



4.1.2 Base Model: Predicting 2016 Beta

In [86]: coefficients_lassob

```
Out [86]:
```

	Percent_male_	under_18_	20_to_40_	40_to_60_	60_above_	median_income_
2006	-0.640029	0.882617	-0.108982	0.444397	-0	-0.0501331
2007	-0.434509	1.01228	-0.0344746	0.643091	-0	-0.379183
2008	-0.661512	0.87094	-0	0.24985	-0	-0.0917136
2009	-0.519636	0.733181	-0.000297521	0.354195	-0	-0.465388
2010	-0.654053	0.317147	-0	0.436378	-0.394344	-0.404777
2011	-0.510295	0.422349	-0.000609473	0.278687	-0	-0.352267
2012	-0.612182	0.443794	-0	0.80993	-0.453153	-0.650738
2013	-0.648021	0	-0	0.378217	-0.586899	-0.611886
2014	-0.722185	-0.802668	-0.196126	0.0961882	-1.42999	-0.733343
2015	-0.837259	-0	0	0.796936	-1.01884	-0.773814

In [87]: coefficients_ridgeb

```
Out [87]:
```

	Percent_male_	under_18_	20_to_40_	40_to_60_	60_above_	median_income_
2006	-0.635164	0.697911	-0.35722	0.346686	-0.313517	-0.124067

2007	-0.434678	0.643388	-0.400622	0.363601	-0.457374	-0.394092
2008	-0.665104	0.663585	-0.309884	0.142336	-0.397116	-0.180151
2009	-0.530811	0.524614	-0.263348	0.229722	-0.345614	-0.514892
2010	-0.663024	0.344533	-0.0204428	0.487956	-0.437737	-0.444311
2011	-0.603805	0.382543	-0.253286	0.420124	-0.422726	-0.560229
2012	-0.617557	0.315643	-0.161369	0.699032	-0.642465	-0.655928
2013	-0.686196	0.125453	0.0942125	0.606812	-0.57452	-0.685282
2014	-0.682826	-0.138663	0.362999	0.512965	-0.629696	-0.687292
2015	-0.853578	0.0652424	0.193498	0.923873	-0.881284	-0.800873

In [88]: *# our x values are the indexes (year 2006 to 2015)*

```
x = coefficients_ridgeb.index.values
```

```
# initialize dataframes to store projected betas for 2016 model
```

```
model_lassob = pd.DataFrame()
```

```
model_lassob['columns'] = base_col
```

```
model_ridgeb = pd.DataFrame()
```

```
model_ridgeb['columns'] = base_col
```

```
# empty lists to fill
```

```
predict_lassob = []
```

```
predict_ridgeb = []
```

```
# for every column
```

```
for col in base_col:
```

```
# fit a regression to the betas of each year
```

```
regress_lassob = LinearRegression().fit(x.reshape(10, 1), coefficients_lassob[col])
```

```
regress_ridgeb = LinearRegression().fit(x.reshape(10, 1), coefficients_ridgeb[col])
```

```
# predict the beta of 2016
```

```
predict_lassob.append(regress_lassob.predict(2016).tolist()[0])
```

```
predict_ridgeb.append(regress_ridgeb.predict(2016).tolist()[0])
```

```
# model the constant
```

```
regress_lassocb = LinearRegression().fit(x.reshape(10,1), constant_lassob)
```

```
const_lassob = regress_lassocb.predict(2016)[0]
```

```
regress_ridgecb = LinearRegression().fit(x.reshape(10,1), constant_ridgeb)
```

```
const_ridgeb = regress_ridgecb.predict(2016)[0]
```

```
# add to df
```

```
model_lassob['coefficients'] = predict_lassob
```

```
model_ridgeb['coefficients'] = predict_ridgeb
```

In [89]: *# show*

```
model_lassob
```

Out [89]: columns coefficients

0	Percent_male_	-0.752476
1	under_18_	-0.470898
2	20_to_40_	-0.039064
3	40_to_60_	0.488650
4	60_above_	-1.157626
5	median_income_	-0.854546

```
In [90]: # show
         model_ridgeb
```

```
Out[90]:
```

	columns	coefficients
0	Percent_male_	-0.770915
1	under_18_	-0.119173
2	20_to_40_	0.301633
3	40_to_60_	0.803401
4	60_above_	-0.779495
5	median_income_	-0.878323

4.1.3 Base Model Performance

```
In [91]: # column of 1
         ones = np.ones((x_yearb[2016].shape[0], 1))

         # add constant to coefficient list
         ylassob = model_lassob['coefficients'].values.tolist()
         ylassob.insert(0, const_lassob)

         # add constant to coefficient list
         yridgeb = model_ridgeb['coefficients'].values.tolist()
         yridgeb.insert(0, const_ridgeb)

         # predict
         lassob_hat = np.dot(np.hstack((ones, x_yearb[2016])), np.array([ylassob]).T)
         ridgeb_hat = np.dot(np.hstack((ones, x_yearb[2016])), np.array([yridgeb]).T)

         # report
         print('The R2 on the test set for using lasso to model the coefficients is', r2_score(y
         print('The R2 on the test set for using ridge to model the coefficients is', r2_score(y
```

The R2 on the test set for using lasso to model the coefficients is 0.0581766303957
The R2 on the test set for using ridge to model the coefficients is 0.0493179030404

```
In [92]: # start a results table
         test_acc = pd.DataFrame(np.c_[r2_score(y_yearb[2016] , lassob_hat.ravel()), r2_score(y_y
```

4.2 Model 1 (Ridge and Lasso)

We use the same procedure we used for the base model but include more features.

```

In [93]: # make dataframe to fill
coefficients_lasso = pd.DataFrame()
coefficients_ridge = pd.DataFrame()

constant_lasso = []
constant_ridge = []

# add a column of column names
coefficients_lasso['columns'] = columns
coefficients_ridge['columns'] = columns

# for every year fit a lasso and ridge model but coef into a df
for i in range(2006, 2016):

    lasso = LassoCV().fit(x_year[i], y_year[i].values.reshape(-1, 1))
    coefficients_lasso[i] = lasso.coef_

    ridge = RidgeCV().fit(x_year[i], y_year[i].values.reshape(-1, 1))
    coefficients_ridge[i] = ridge.coef_[0]

# constants
constant_lasso.append(lasso.intercept_)
constant_ridge.append(ridge.intercept_)

In [94]: coefficients_lasso = rename(coefficients_lasso.T)
coefficients_ridge = rename(coefficients_ridge.T)

In [95]: coefficients_lasso

Out[95]:
    Percent_male_ under_18_ 20_to_40_ 40_to_60_ 60_above_ Now married_ \
2006 -0.0737585          0 -0.120371         -0          0 -0.310646
2007  0.257747          0 -0.159392         -0  0.0311602 -0.0886804
2008          -0  0.199566 -0.177468         -0          0 -0.249463
2009          0          0          -0          0         -0          -0
2010          -0          0 -0.170635          0         -0 -0.200519
2011          0          0 -0.288664          0          0          -0
2012          -0          0 -0.288608          0         -0 -0.170833
2013          -0  0.311316 -0.20135          0         -0 -0.33857
2014          -0          0          -0  0.0918426 -0.111397 -0.285857
2015 -0.140188          0          -0          0 -0.121281 -0.0967096

    Widowed_ Divorced_ Separated_ Never married_ ... poverty_ \
2006 0.0783093 0.260305          -0          0 ... 0.00942778
2007 0.421492 0.180766 0.265233          0 ... 0.0635628
2008 0.178018 0.274448          0          0 ...          -0
2009 0.301115 0.20434          0          0 ... 0.109068
2010          0 0.0416709          0          0 ...          0

```

2011	0	0.190487	0	0	...	0.0897766
2012	0	0.0325993	0	0	...	0.111312
2013	0	0.40593	-0.245987	0	...	0.0919142
2014	0	0.197677	0	0	...	0
2015	0	0.0941221	-0	0	...	0

	median_income_	White alone_Percent_ \
2006	-0	-0.736414
2007	-0.18895	-1.12757
2008	-0	-1.04499
2009	-0.0271648	-0.978372
2010	-0	-0.670061
2011	-0.109286	-1.07249
2012	-0.000290382	-1.16685
2013	-0	-0.808767
2014	-0.04932	-0.610424
2015	-0	-0.898325

	Black or African American alone_Percent_ \
2006	1.30109
2007	1.37413
2008	1.06773
2009	0.810164
2010	0.893656
2011	0.889942
2012	0.774943
2013	0.871078
2014	1.0114
2015	1.03636

	American Indian and Alaska Native alone_Percent_	Asian alone_Percent_ \
2006	-0.0458542	-0
2007	0.268435	-0
2008	-0	-0
2009	-0	-0
2010	-0	0
2011	-0	-0
2012	-0	0
2013	0	0.021704
2014	-0	-0
2015	0	-0.0358721

	Native Hawaiian and Other Pacific Islander alone_Percent_ \
2006	0
2007	-0
2008	-0.0584644
2009	-0
2010	-0

2011	0
2012	-0
2013	-0.0942926
2014	-0.0464559
2015	-0.0821578

	Some other race alone_Percent_	Two or more races:_Percent_	gdp_
2006	0.292345	-0	0.418159
2007	0	-0	0.297181
2008	0	-0	0.0381334
2009	0	0	-0
2010	0	0.0586398	0
2011	0	0	0
2012	0	0	0
2013	-0	-0	0
2014	0	-0.000114191	0
2015	0	0	0

[10 rows x 22 columns]

In [96]: coefficients_ridge

```
Out[96]:
```

	Percent_male_	under_18_	20_to_40_	40_to_60_	60_above_	\
2006	-0.0747912	0.238735	-0.287168	-0.104337	0.147873	
2007	0.403424	0.144366	-0.374893	-0.144828	0.204242	
2008	0.0231769	0.405794	-0.350136	-0.140072	0.0393436	
2009	0.153332	0.167023	-0.137943	0.0373835	-0.0340698	
2010	0.098095	0.388321	-0.365263	0.15625	-0.0391539	
2011	0.207791	0.313842	-0.639854	0.00415628	0.225246	
2012	0.144069	0.468521	-0.647318	0.129458	0.00624488	
2013	-0.0428088	0.420285	-0.401608	-0.0114577	-0.161821	
2014	-0.0604756	0.0846521	-0.00468646	0.257999	-0.359619	
2015	-0.433792	0.256249	0.0241171	0.141462	-0.373812	

	Now married_	Widowed_	Divorced_	Separated_	Never married_	...	\
2006	-0.371726	0.135581	0.420494	-0.180483	0.21671	...	
2007	-0.217077	0.357562	0.221052	0.339568	0.0142137	...	
2008	-0.353166	0.284775	0.389001	-0.0700646	0.122847	...	
2009	-0.0807763	0.489216	0.266589	-0.223994	0.0287585	...	
2010	-0.245115	0.261595	0.103557	-0.0503443	0.119954	...	
2011	-0.160869	-0.252273	0.234301	-0.164846	0.111794	...	
2012	-0.276079	0.011981	0.173533	-0.184547	0.204041	...	
2013	-0.397335	0.0842268	0.550506	-0.521681	0.15544	...	
2014	-0.294289	0.306807	0.402417	-0.0973602	0.0369984	...	
2015	-0.307489	0.312615	0.439992	-0.446619	0.103359	...	

	poverty_	median_income_	White alone_Percent_	\
2006	0.0663278	-0.038225	-1.00301	

2007	0.0646792	-0.244484	-1.18221
2008	-0.305116	-0.208007	-1.04532
2009	0.197809	-0.146761	-0.979988
2010	0.0890927	-0.17026	-0.724182
2011	0.0383113	-0.427316	-1.04477
2012	0.0856101	-0.398206	-1.00022
2013	0.107768	-0.0111939	-0.813019
2014	-0.0426752	-0.162488	-0.777346
2015	-0.452774	-0.196361	-0.958669

	Black or African American alone_Percent_ \
2006	1.03808
2007	1.26932
2008	1.07242
2009	0.988476
2010	0.839109
2011	1.12364
2012	1.01552
2013	0.858185
2014	0.831626
2015	0.989301

	American Indian and Alaska Native alone_Percent_ Asian alone_Percent_ \	
2006	-0.239006	-0.127095
2007	0.278785	-0.0687487
2008	-0.0468554	-0.0444688
2009	-0.115869	-0.140838
2010	-0.119499	0.00278037
2011	-0.107363	0.0475343
2012	-0.114439	0.176301
2013	0.27857	0.165308
2014	-0.1917	-0.0529405
2015	0.143204	-0.244905

	Native Hawaiian and Other Pacific Islander alone_Percent_ \
2006	0.0935778
2007	0.00269123
2008	-0.140197
2009	-0.057806
2010	-0.0341539
2011	0.0738702
2012	-0.0599582
2013	-0.0932266
2014	-0.042738
2015	-0.0791229

	Some other race alone_Percent_ Two or more races:_Percent_ gdp_
2006	0.274821 -0.112496 0.662554


2007	-0.13859	0.00955613	0.434763
2008	0.128263	0.0255226	0.182589
2009	0.208286	0.190136	0.136891
2010	-0.134649	0.254433	0.167132
2011	-0.00728836	0.0476045	0.197833
2012	0.0336246	0.15784	0.173905
2013	-0.201819	-0.0851694	0.130449
2014	0.130912	-0.0468135	0.128606
2015	0.0926911	0.124646	0.294837

[10 rows x 22 columns]

4.2.1 Lasso Beta Visualization

In [97]: `plt.scatter(np.arange(2006, 2016, 1), constant_lasso)`

Out [97]: `<matplotlib.collections.PathCollection at 0x12531d550>`



In [98]: `sns.lmplot(x = 'index', y = 'Percent_male_', data= coefficients_lasso.reset_index(), fi`

Out [98]: `<seaborn.axisgrid.FacetGrid at 0x1246b7780>`



```
In [99]: sns.lmplot(x = 'index', y = "median_income_", data= coefficients_lasso.reset_index(), f
```

```
Out[99]: <seaborn.axisgrid.FacetGrid at 0x1242b5be0>
```



```
In [100]: sns.lmplot(x = 'index', y = "White alone_Percent_", data= coefficients_lasso.reset_ind
```

```
Out[100]: <seaborn.axisgrid.FacetGrid at 0x1252f2198>
```




```
In [101]: sns.lmplot(x = 'index', y = "gdp_", data= coefficients_lasso.reset_index(), fit_reg =  
Out[101]: <seaborn.axisgrid.FacetGrid at 0x11f568518>
```



4.2.2 Ridge Beta Visualization

```
In [102]: plt.scatter(np.arange(2006, 2016, 1), constant_ridge)
```

```
Out[102]: <matplotlib.collections.PathCollection at 0x12004d7f0>
```



```
In [103]: sns.lmplot(x = 'index', y = "Percent bachelor's degree or higher_", data= coefficients
```

```
Out[103]: <seaborn.axisgrid.FacetGrid at 0x120d44320>
```




```
In [104]: sns.lmplot(x = 'index', y = "median_income_", data= coefficients_ridge.reset_index(),
```

```
Out[104]: <seaborn.axisgrid.FacetGrid at 0x11d065e80>
```




```
In [105]: sns.lmplot(x = 'index', y = "White alone_Percent_", data= coefficients_ridge.reset_ind
```

```
Out[105]: <seaborn.axisgrid.FacetGrid at 0x120abdc50>
```



```
In [106]: sns.lmplot(x = 'index', y = "gdp_", data= coefficients_ridge.reset_index(), fit_reg =  
Out[106]: <seaborn.axisgrid.FacetGrid at 0x11f9be8d0>
```



4.3 Predicting 2016 Beta

```
In [107]: # initialize dataframes to store projected betas for 2016 model
select_col = []
select_col_16 = []

model_lasso = pd.DataFrame()
model_ridge = pd.DataFrame()

# empty lists to fill
predict_lasso = []
predict_ridge = []

# for every column
for col in columns:
    if len(coefficients_lasso[col][coefficients_lasso[col].abs() > 0]) > 4:

        select_col.append(col)
        select_col_16.append(col + str(16))
```

```

# fit a regression to the betas of each year
regress_lasso = LinearRegression().fit(x.reshape(10, 1), coefficients_lasso[co
regress_ridge = LinearRegression().fit(x.reshape(10, 1), coefficients_ridge[co

# predict the beta of 2016
predict_lasso.append(regress_lasso.predict(2016).tolist()[0])
predict_ridge.append(regress_ridge.predict(2016).tolist()[0])

# add columns
model_lasso['columns'] = select_col
model_ridge['columns'] = select_col

# model the constant
regress_lassoc = LinearRegression().fit(x.reshape(10,1), constant_lasso)
const_lasso = regress_lassoc.predict(2016)[0]
regress_ridgec = LinearRegression().fit(x.reshape(10,1), constant_ridge)
const_ridge = regress_ridgec.predict(2016)[0]

# add to df
model_lasso['coefficients'] = predict_lasso
model_ridge['coefficients'] = predict_ridge

```

In [108]: model_lasso

```

Out[108]:
           columns  coefficients
0          20_to_40_    -0.104122
1      Now married_    -0.181205
2        Divorced_     0.152026
3  Percent high school graduate or higher_    -0.071966
4    Percent bachelor's degree or higher_    -0.399958
5                poverty_     0.048383
6          median_income_    -0.005876
7      White alone_Percent_    -0.832224
8  Black or African American alone_Percent_     0.802572

```

In [109]: model_ridge

```

Out[109]:
           columns  coefficients
0          20_to_40_    -0.207377
1      Now married_    -0.293221
2        Divorced_     0.390282
3  Percent high school graduate or higher_    -0.365598
4    Percent bachelor's degree or higher_    -0.379171
5                poverty_    -0.139975
6          median_income_    -0.229550
7      White alone_Percent_    -0.819096
8  Black or African American alone_Percent_     0.862290

```

```

In [110]: ones = np.ones((x_year[2016].shape[0],1))

# add constant to coefficient list
ylasso = model_lasso['coefficients'].values.tolist()
ylasso.insert(0, const_lasso)

# add constant to coefficient list
yridge = model_ridge['coefficients'].values.tolist()
yridge.insert(0, const_ridge)

# predict
lasso_hat = np.dot(np.hstack((ones, x_year[2016][select_col_16])), np.array([ylasso])).
ridge_hat = np.dot(np.hstack((ones, x_year[2016][select_col_16])), np.array([yridge])).

# report
print('The R2 on the test set for using lasso to model the coefficients is', r2_score(
print('The R2 on the test set for using ridge to model the coefficients is', r2_score(

```

The R2 on the test set for using lasso to model the coefficients is 0.405715521583
The R2 on the test set for using ridge to model the coefficients is 0.444669689189

```

In [111]: test_acc['Lasso (base + extra features)'] = r2_score(y_year[2016] , lasso_hat.ravel())
test_acc['Ridge (base + extra features)'] = r2_score(y_year[2016] , ridge_hat.ravel())

In [112]: test_acc

```

```

Out[112]:      Lasso (Base)  Ridge (Base)  Lasso (base + extra features)  \
0      0.058177      0.049318      0.405716

      Ridge (base + extra features)
0      0.44467

```

4.4 Average

In this method, we average over all years and take a subset of the averaged dataframe to train and test

```

In [113]: colnames = dict_df_year[2006].columns.str.replace('06', '').tolist()

In [114]: # msa_df has all the msa names, good place to start search
msa_list = msa_df['msa_name'].values
df_avg = pd.DataFrame(columns = colnames)

In [115]: # iterate over msas, look at trends in years for each feature
# average over these feature-trends
for i in range(len(msa_list)):
    new_row = []
    for j in range((dict_avg[2006].shape[1])):

```

```

var_list = []
for k in range(10):
    if dict_avg[2006 + k].loc[dict_avg[2006 + k]['msa_name'] == msa_list[i]].e
        var_list.append(dict_avg[2006 + k].loc[dict_avg[2006 + k]['msa_name']
    else:
        var_list.append(np.nan)
var_list = np.array(var_list)
if j == 20:
    avg = msa_list[i]
elif j == 22:
    avg = unique_id.loc[unique_id['msa_name'] == msa_list[i]]['ID'].values[0]
else:
    if i == 3 and j == 0:
        print((var_list))
        print(dict_avg[2006 + k].loc[dict_avg[2006 + k]['msa_name'] == msa_list[i]]['ID'].values[0])
    avg = np.mean(var_list[~np.isnan(var_list)])
new_row.append(avg)
df_avg.loc[i] = new_row

[ 0.47601507  0.48071169          nan  0.47601042  0.477668      0.47350217
 0.47447732  0.47534826  0.46781698  0.46696434]
Percent_male_15  Percent_female_15  AGE - 80 to 84 years_15  \
2          0.466964          0.533036          1.9

AGE - 85 years and over_15  Now married_15  Widowed_15  Divorced_15  \
2          1.6          39.4          7.8          12.6

Separated_15  Never married_15  Percent high school graduate or higher_15  \
2          5.2          35.0          82.3

...          Two or more races:_Percent_15  msa_name  Murder 2015  ID  \
2          ...          0.017504  Albany, GA          12.0  10500

gdp_15  murder_rate  under_18_15  20_to_40_15  40_to_60_15  60_above_15
2  29235      7.643458          28.5          24.4          26.3          20.9

[1 rows x 29 columns]

In [116]: df_avg = df_avg.dropna(axis = 0)

In [117]: base_col = ['Percent_male_', 'under_18_', '20_to_40_', '40_to_60_', '60_above_', 'med

# train/test split
np.random.seed(9001)
msk = np.random.randn(df_avg.shape[0]) < 0.7
df_train = df_avg[msk]
df_test = df_avg[~msk]

```

```

# specific training/testing columns
x_train = df_train[base_col]
x_test = df_test[base_col]
y_train = df_train['murder_rate']
y_test = df_test['murder_rate']

# standardizing
scaler = StandardScaler().fit(x_train)
x_train = scaler.transform(x_train)
x_test = scaler.transform(x_test)

```

```

In [118]: lassoavg = LassoCV(fit_intercept = True).fit(x_train, y_train)
print('Test R^2 with Lasso = %s'%(lassoavg.score(x_test, y_test)))
ridgeavg = RidgeCV(fit_intercept = True).fit(x_train, y_train)
print('Test R^2 with Ridge = %s'%(ridgeavg.score(x_test, y_test)))

```

```

Test R^2 with Lasso = 0.152184287169
Test R^2 with Ridge = 0.155698225619

```

```

In [119]: test_acc['Lasso (Base Average)'] = lassoavg.score(x_test, y_test)
test_acc['Ridge (Base Average)'] = ridgeavg.score(x_test, y_test)

```

```

In [120]: plt.figure(figsize=(15,9))
ypred = lassoavg.predict(x_test)[0:40]
plt.scatter(y_test[0:40], ypred, c='purple', marker = 'x', label = 'Test')
plt.scatter(y_train[0:40], lassoavg.predict(x_train)[0:40], c = 'orange', marker = '^')

xvar = np.linspace(0,120,100)
plt.text(10,3,'Test R^2 = ' + str(round(lassoavg.score(x_test,y_test),3)),fontsize=24)
yvar = xvar
plt.plot(xvar,yvar,'--',c='g')
plt.xlabel('Test', fontsize = 24)
plt.ylabel('Predicted',fontsize = 24)
plt.legend(loc=2,prop={'size':24})


plt.ylim(0, 20)
plt.xlim(0, 20)

```

```

Out[120]: (0, 20)

```

```
In [121]: base_col = ['Percent_male_', 'Now married_', 'Widowed_', 'Divorced_', 'Separated_', 'Ne
```

```
In [122]: # train/test split
np.random.seed(9001)
msk = np.random.randn(df_avg.shape[0]) < 0.7
df_train = df_avg[msk]
df_test = df_avg[~msk]

# specific training/testing columns
x_train = df_train[base_col]
x_test = df_test[base_col]
y_train = df_train['murder_rate']
y_test = df_test['murder_rate']

# standardizing
scaler = StandardScaler().fit(x_train)
x_train = scaler.transform(x_train)
x_test = scaler.transform(x_test)

lassoavg = LassoCV(fit_intercept = True).fit(x_train, y_train)
print('Test R^2 with Lasso = %s'%(lassoavg.score(x_test, y_test)))
ridgeavg = RidgeCV(fit_intercept = True).fit(x_train, y_train)
print('Test R^2 with Ridge = %s'%(ridgeavg.score(x_test, y_test)))
rfavg = RandomForestRegressor().fit(x_train, y_train)
print('Test R^2 with RF = %s'%(rfavg.score(x_test, y_test)))
```

Test R^2 with Lasso = 0.604330356174
Test R^2 with Ridge = 0.639956016422
Test R^2 with RF = 0.595157101937

```
In [123]: importances = rfavg.feature_importances_
std = np.std([tree.feature_importances_ for tree in rfavg.estimators_],
              axis=0)
indices = np.argsort(importances)[::-1]

newlab = [base_col[i] for i in indices]

# Print the feature ranking
print("Feature ranking:")


for f in range(x_train.shape[1]):
    print("%d. feature %d (%f)" % (f + 1, indices[f], importances[indices[f]]))

# Plot the feature importances of the forest
plt.figure(figsize=(15,9))
plt.title("Feature importances")
plt.bar(range(x_train.shape[1]), importances[indices],
        color="r", yerr=std[indices], align="center")
plt.xticks(range(x_train.shape[1]), newlab,rotation=90)
plt.xlim([-1, x_train.shape[1]])
plt.show()
```

Feature ranking:

1. feature 11 (0.388603)
2. feature 10 (0.162002)
3. feature 4 (0.089568)
4. feature 16 (0.032254)
5. feature 0 (0.031671)
6. feature 3 (0.027373)
7. feature 1 (0.026191)
8. feature 18 (0.025638)
9. feature 15 (0.020576)
10. feature 12 (0.020480)
11. feature 5 (0.019303)
12. feature 6 (0.019169)
13. feature 14 (0.018102)
14. feature 17 (0.017262)
15. feature 7 (0.017215)
16. feature 21 (0.014445)
17. feature 20 (0.013815)
18. feature 19 (0.013814)
19. feature 9 (0.013720)
20. feature 2 (0.011951)

21. feature 13 (0.011079)
 22. feature 8 (0.005770)



```
In [124]: test_acc['Lasso (base+add features average)'] = lassoavg.score(x_test, y_test)
          test_acc['Ridge (base+add features average)'] = ridgeavg.score(x_test, y_test)
          test_acc['RF (base+add features average)'] = rfavg.score(x_test, y_test)
```

```
In [125]: plt.figure(figsize=(15, 9))
          ypred = lassoavg.predict(x_test)[0:40]
          plt.scatter(y_test[0:40], ypred, c='purple', marker = 'x', label = 'Test')
          plt.scatter(y_train[0:40], lassoavg.predict(x_train)[0:40], c='orange', marker = '^', label = 'Train')
```


```

xvar = np.linspace(0,20 ,10)
plt.text(10, 3, 'Test R^2 = ' + str(round(lassoavg.score(x_test, y_test), 3)),fontSize=24)
yvar = xvar
plt.plot(xvar,yvar,'--',c='g')
plt.xlabel('Test',fontSize=24)
plt.ylabel('Predicted',fontSize=24)
plt.legend(loc = 2,prop = {'size':24})

plt.ylim(0, 20)
plt.xlim(0, 20)

```

Out[125]: (0, 20)



```

In [126]: # train/test split
np.random.seed(9001)
msk = np.random.randn(df_avg.shape[0]) < 0.7
df_train = df_avg[msk]
df_test = df_avg[~msk]

# specific training/testing columns
x_train = df_train[base_col]
x_test = df_test[base_col]
y_train = df_train['murder_rate']
y_test = df_test['murder_rate']

```

```

# standardizing
scaler = StandardScaler().fit(x_train)
x_train = scaler.transform(x_train)
x_test = scaler.transform(x_test)

from sklearn.preprocessing import PolynomialFeatures

x_train, x_test = more_terms(x_train, x_test, 3)

lassoavg = LassoCV(fit_intercept = True).fit(x_train, y_train)
print('Test R^2 with Lasso = %s'%(lassoavg.score(x_test, y_test)))
ridgeavg = RidgeCV(fit_intercept = True).fit(x_train, y_train)
print('Test R^2 with Ridge = %s'%(ridgeavg.score(x_test, y_test)))
rfavg = RandomForestRegressor().fit(x_train, y_train)
print('Test R^2 with RF = %s'%(rfavg.score(x_test, y_test)))

```

Test R^2 with Lasso = 0.456706649263

Test R^2 with Ridge = -22.021980383

Test R^2 with RF = 0.565370070829

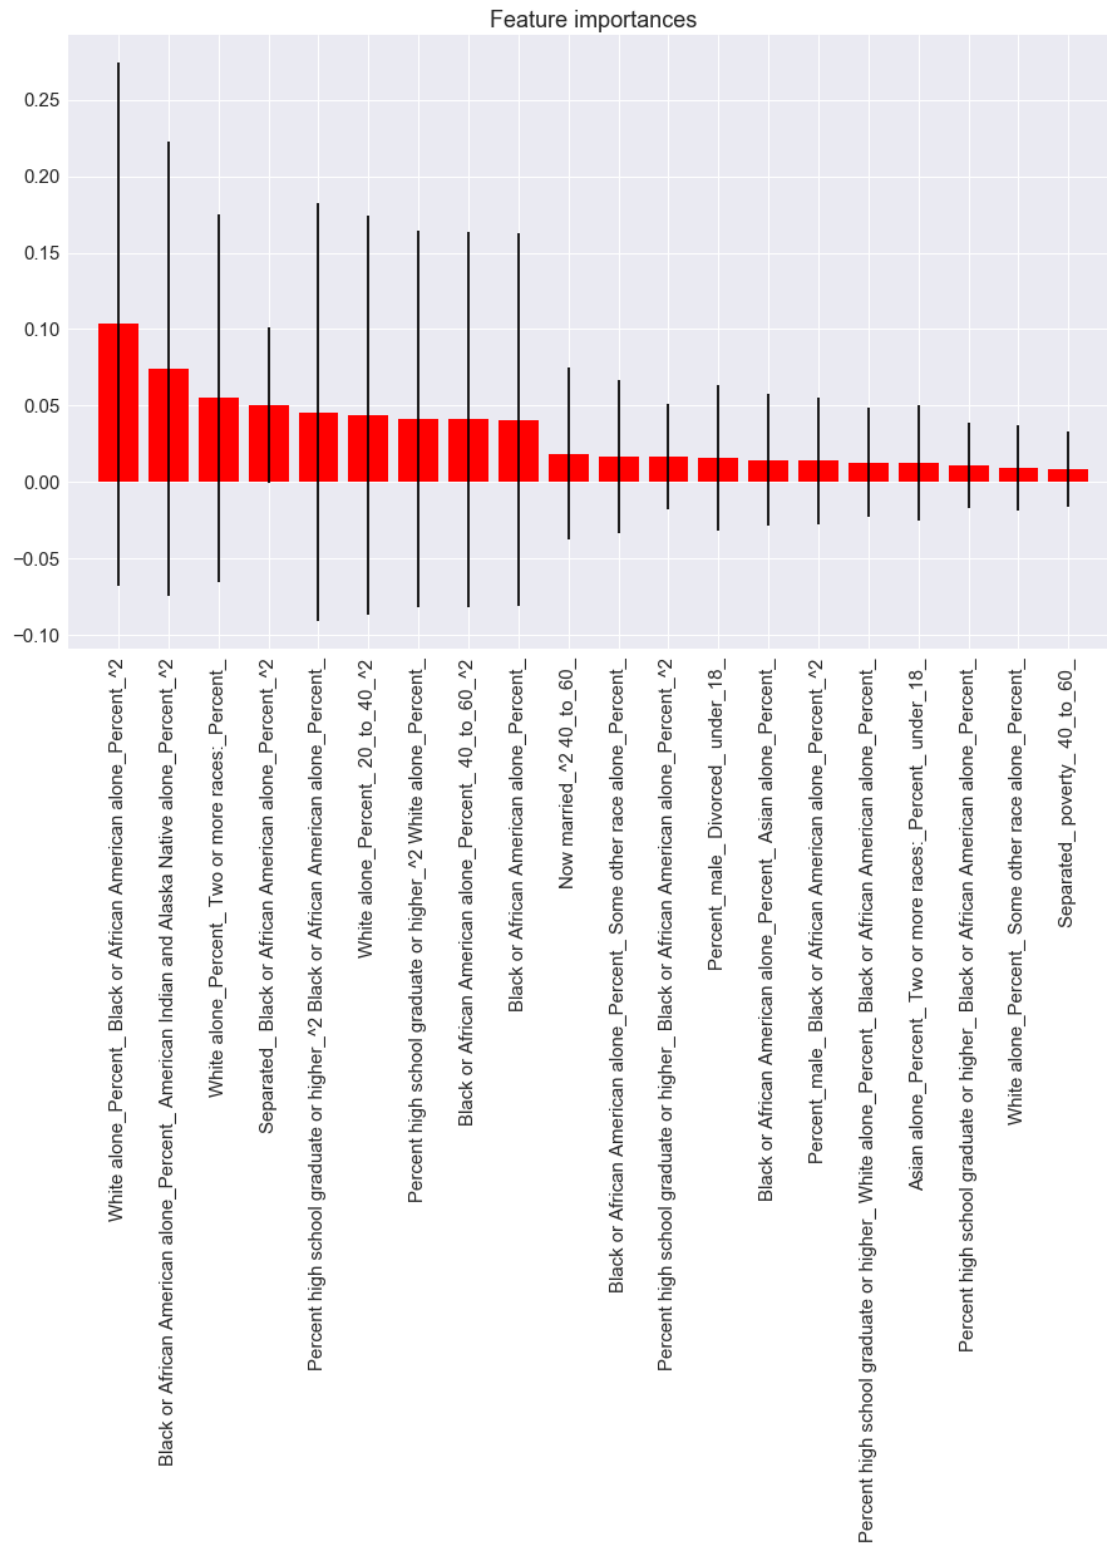
```

In [127]: importances = rfavg.feature_importances_
std = np.std([tree.feature_importances_ for tree in rfavg.estimators_],
              axis=0)
indices = np.argsort(importances)[::-1]

newlab = [x_train.columns[i] for i in indices]

# Plot the feature importances of the forest
plt.figure(figsize=(15,9))
plt.title("Feature importances")
plt.bar(range(20), importances[indices][0:20],
        color="r", yerr=std[indices][0:20], align="center")
plt.xticks(range(20), newlab[0:20],rotation=90)
plt.xlim([-1, 20])
plt.show()

```



```
In [128]: test_acc['Lasso (complex features average)'] = lassoavg.score(x_test, y_test)
```

```
test_acc['Ridge (complex features average)'] = ridgeavg.score(x_test, y_test)
test_acc['RF (complex features average)'] = rfavg.score(x_test, y_test)
```

```
In [129]: # train/test split
np.random.seed(9001)
msk = np.random.randn(df_avg.shape[0]) < 0.7
df_train = df_avg[msk]
df_test = df_avg[~msk]


# specific training/testing columns
x_train = df_train[base_col]
x_test = df_test[base_col]
y_train = df_train['murder_rate']
y_test = df_test['murder_rate']

# standardizing
scaler = MinMaxScaler().fit(x_train)
x_train = scaler.transform(x_train)
x_test = scaler.transform(x_test)

pcavar = []
i = 1
while True:
    pca = PCA(n_components = i)
    pca.fit(x_train)
    pcavar.append(pca.explained_variance_ratio_.sum())
    if (pca.explained_variance_ratio_.sum()) >= 0.9:
        break
    i += 1

plt.figure(figsize = (10,8))
plt.plot(pcavar)
plt.title('# PCA components vs total variance')
plt.ylabel('Variance')
plt.xlabel('# of PCA components')
```

```
Out[129]: <matplotlib.text.Text at 0x12375c748>
```



```
In [130]: pca = PCA(n_components = 7)
pca.fit(x_train)
x_train = pca.transform(x_train)
x_test = pca.transform(x_test)

lassoavg = LassoCV(fit_intercept = True).fit(x_train, y_train)
print('Test R^2 with Lasso = %s'%(lassoavg.score(x_test, y_test)))
ridgeavg = RidgeCV(fit_intercept = True).fit(x_train, y_train)
print('Test R^2 with Ridge = %s'%(ridgeavg.score(x_test, y_test)))
rfavg = RandomForestRegressor().fit(x_train, y_train)
print('Test R^2 with RF = %s'%(rfavg.score(x_test, y_test)))
```

```
Test R^2 with Lasso = 0.57626498613
Test R^2 with Ridge = 0.577600708353
Test R^2 with RF = 0.529015422912
```

```
In [131]: test_acc['Lasso (PCA)'] = lassoavg.score(x_test, y_test)
test_acc['Ridge (PCA)'] = ridgeavg.score(x_test, y_test)
test_acc['RF (PCA)'] = rfavg.score(x_test, y_test)
```


5 2015 model

In this approach, we only use the 2015 dataset as train and the 2016 as test.

```
In [132]: base_col = ['Percent_male_', 'under_18_', '20_to_40_', '40_to_60_', '60_above_', 'med
dict15 = pd.DataFrame.copy(dict_df_year[2015])
colnames = dict15.columns.str.replace('15', '').tolist()
dict15.columns = colnames
dict15 = dict15.dropna(axis = 0)

dict16 = pd.DataFrame.copy(dict_df_year[2016])
colnames = dict16.columns.str.replace('16', '').tolist()
dict16.columns = colnames
dict16 = dict16.dropna(axis = 0)

In [133]: x_train = dict15[base_col]
y_train = dict15['murder_rate']
x_test = dict16[base_col]
y_test = dict16['murder_rate']

scaler = StandardScaler().fit(x_train)
x_train = scaler.transform(x_train)
x_test = scaler.transform(x_test)

lassoavg = LassoCV(fit_intercept = True).fit(x_train, y_train)
print('Test R^2 with Lasso = %s'%(lassoavg.score(x_test, y_test)))
ridgeavg = RidgeCV(fit_intercept = True).fit(x_train, y_train)
print('Test R^2 with Ridge = %s'%(ridgeavg.score(x_test, y_test)))
rfavg = RandomForestRegressor().fit(x_train, y_train)
print('Test R^2 with RF = %s'%(rfavg.score(x_test, y_test)))

Test R^2 with Lasso = 0.184984007671
Test R^2 with Ridge = 0.185231612241
Test R^2 with RF = 0.382931023462

In [134]: test_acc['Lasso (base 2015)'] = lassoavg.score(x_test, y_test)
test_acc['Ridge (base 2015)'] = ridgeavg.score(x_test, y_test)
test_acc['RF (base 2015)'] = rfavg.score(x_test, y_test)

In [135]: base_col = ['Percent_male_', 'Now married_', 'Widowed_', 'Divorced_', 'Separated_', 'Ne
x_train = dict15[base_col]
y_train = dict15['murder_rate']
x_test = dict16[base_col]
y_test = dict16['murder_rate']

scaler = StandardScaler().fit(x_train)
x_train = scaler.transform(x_train)
```

```

x_test = scaler.transform(x_test)

lassoavg = LassoCV(fit_intercept = True).fit(x_train, y_train)
print('Test R^2 with Lasso = %s'%(lassoavg.score(x_test, y_test)))
ridgeavg = RidgeCV(fit_intercept = True).fit(x_train, y_train)
print('Test R^2 with Ridge = %s'%(ridgeavg.score(x_test, y_test)))
rfavg = RandomForestRegressor().fit(x_train, y_train)
print('Test R^2 with RF = %s'%(rfavg.score(x_test, y_test)))

```

```

Test R^2 with Lasso = 0.539690336936
Test R^2 with Ridge = 0.542384805987
Test R^2 with RF = 0.563916870639

```

```

In [136]: test_acc['Lasso (base+add 2015)'] = lassoavg.score(x_test, y_test)
test_acc['Ridge (base+add 2015)'] = ridgeavg.score(x_test, y_test)
test_acc['RF (base+add 2015)'] = rfavg.score(x_test, y_test)

```

```

In [137]: x_train = dict15[base_col]
y_train = dict15['murder_rate']
x_test = dict16[base_col]
y_test = dict16['murder_rate']

scaler = StandardScaler().fit(x_train)
x_train = scaler.transform(x_train)
x_test = scaler.transform(x_test)

x_train, x_test = more_terms(x_train, x_test, 3)

lassoavg = LassoCV(fit_intercept = True).fit(x_train, y_train)
print('Test R^2 with Lasso = %s'%(lassoavg.score(x_test, y_test)))
ridgeavg = RidgeCV(fit_intercept = True).fit(x_train, y_train)
print('Test R^2 with Ridge = %s'%(ridgeavg.score(x_test, y_test)))
rfavg = RandomForestRegressor().fit(x_train, y_train)
print('Test R^2 with RF = %s'%(rfavg.score(x_test, y_test)))

```

```

Test R^2 with Lasso = 0.429372422981
Test R^2 with Ridge = -0.247138270863
Test R^2 with RF = 0.506115802568

```

```

In [138]: test_acc['Lasso (complex 2015)'] = lassoavg.score(x_test, y_test)
test_acc['Ridge (complex 2015)'] = ridgeavg.score(x_test, y_test)
test_acc['RF (complex 2015)'] = rfavg.score(x_test, y_test)

```

```

In [139]: test_acc.index = ['Test Accuracy']

```

```

In [140]: # summary table
test_acc

```

```

Out[140]:
Test Accuracy      Lasso (Base)      Ridge (Base)      Lasso (base + extra features) \
                                0.058177      0.049318                                0.405716

Test Accuracy      Ridge (base + extra features)      Lasso (Base Average) \
                                0.44467                                0.152184

Test Accuracy      Ridge (Base Average)      Lasso (base+add features average) \
                                0.155698                                0.60433

Test Accuracy      Ridge (base+add features average) \
                                0.639956

Test Accuracy      RF (base+add features average) \
                                0.595157

Test Accuracy      Lasso (complex features average)      ...      RF (PCA) \
                                0.456707      ...      0.529015

Test Accuracy      Lasso (base 2015)      Ridge (base 2015)      RF (base 2015) \
                                0.184984      0.185232      0.382931

Test Accuracy      Lasso (base+add 2015)      Ridge (base+add 2015) \
                                0.53969      0.542385

Test Accuracy      RF (base+add 2015)      Lasso (complex 2015)      Ridge (complex 2015) \
                                0.563917      0.429372      -0.247138

Test Accuracy      RF (complex 2015)
                                0.506116

[1 rows x 24 columns]

```